

M2 - Optimization and Inference Techniques for CV

Final Presentation

Graph Cuts



Group 8

Alex Carrillo
Guillem Martínez

MSc in Computer Vision

Optimization

Criteria(s) examples

“Route length” → solution → Shortest route

“Travel time” → solution → Fastest route

Concept

In **optimization**, we have to find a **solution** for a given problem which is “**best**” in the sense of a **criteria**.

Goal and approach

1. Define a **suitable criteria**
2. **Select best solution** based on an optimization **algorithm**
(different criteria might lead to different “optimal” solutions)

Ingredients

- A sense of the **set of all possible solutions** for our optimization problem
- Something that enables us to calculate a **quantity** which indicates the “**goodness**” of a particular solution
- (optional) **Additional constraints** which the solution has to satisfy

Optimization

Mathematically, the ingredients are:

- a solution set **S**
- an (objective) function **f**

Find the solution **x*** that minimizes (max.) the objective function **f**.

Moreover,

- An optimization **algorithm** has to find a solution in **reasonable time** and **accuracy**
- Ideally, the problem task has to be a **well-posed** problem:
 1. A **solution exists**
 2. The solution is **unique**
 3. Small variations of **initial conditions** result in only **small variations** of the solution
 4. Sometimes, it can be **solved analytically**

Else, we deal with complex **ill-posed problems** (typical in computer vision)

Expressed mathematically:

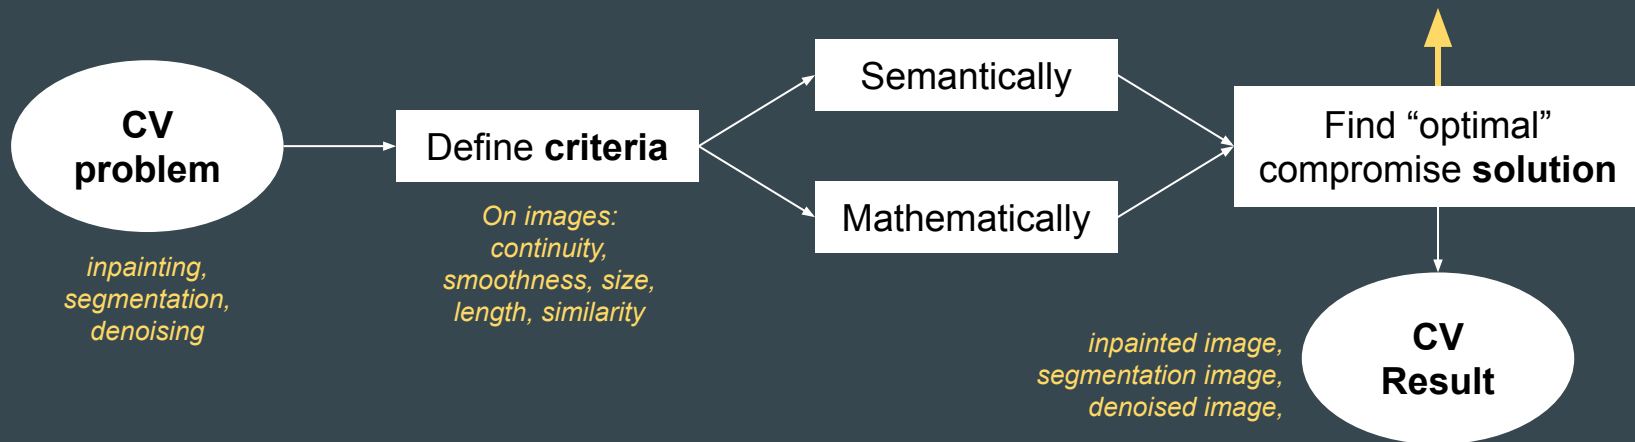
$$f : S \rightarrow \mathbb{R}$$

$$x^* = \arg \min_{x \in S} f(x)$$

Optimization in Computer Vision

Many **optimization problems** arise from **computer vision**:

- Image restoration (denoising, inpainting), segmentation, compression, object detection, 3D (or multi-view) scene reconstruction, optical flow...
- Often, find **acceptable solutions** to an **ill-posed** problem
- **Spatial** and **visual relationships** play an important role in optimization



Chosen Paper: *Interactive Graph Cuts* (Segmentation)

Problem definition

- **Segmentation problem**

- Finding different **regions** or **segments** (set of pixels) that partitionate an image into **meaningful parts**
- **Grouping objects** by **some criteria**, such that those within a group respond similarly and those in a different group respond differently

Goal


- Simplify the representation of an image
- Divide an image into **two segments**: “**object**” and “**background**”

Result

- An **image** that **maps** the different regions or segments (set of pixels)

Approaches

- Fully automatic segmentation
 - (which seems to) Never be perfect...
- **Interactive segmentation**
 - (is evaluated) More reliable...



original image segmentation result




Interactive segmentation (object and background)

Chosen Paper: Interactive Segmentation Criteria

1. User imposes certain **hard constraints** for segmentation
 - 1.1. Indicate certain pixels (*seeds*) that absolutely have to be part of the **object**
 - 1.2. Indicate certain pixels (*seeds*) that absolutely have to be part of the **background**
2. The rest of the image is **segmented automatically**
(globally optimizing among all segmentations satisfying the hard constraints)
3. Get segmentation **results quickly** via very intuitive **interactions**
(when user adds/removes any hard constraints *seeds*)

This (mathematically thinking) translates into:

- Defining a **cost** (function) in terms of **boundary** and **regions properties** of the segments
- The **regions properties** can be viewed as **soft constraints** for segmentation
- **Efficiently** (re)compute a **global optimal segmentation**



*Interactive segmentation
(object and background)*

Chosen Paper: Graph Cuts Segmentation Mathematically

1. Represent image as a **Markov Random Field** represented by a **graph**:

$$\mathcal{V} = \mathcal{P} \cup \{S, T\} \quad \mathcal{E} = \mathcal{N} \bigcup_{p \in \mathcal{P}} \{\{p, S\}, \{p, T\}\}$$

\mathcal{P} = pixel node

S = terminal (source) node / object

T = terminal (sink) node / background

\mathcal{N} = n-links (neighborhood links)



$\{p, S \text{ or } T\}$ = t-links (terminal links)

2. Assign a non-negative weight w_e to each edge \mathcal{E}
3. Define the **segmentation** as a “graph cut” \mathcal{C}
(a set of edges that separate the terminals on the graph)

$$|\mathcal{C}| = \sum_{e \in \mathcal{C}} w_e$$

4. Conditions for a **feasible cut**

- 4.1. \mathcal{C} cuts exactly one t-link at each p
- 4.2. $\{p, q\}$ in \mathcal{C} iff p, q are linked to different terminals
- 4.3. If p in **object**, then $\{p, T\}$ in \mathcal{C}
- 4.4. If p in **background**, then $\{p, S\}$ in \mathcal{C}



Chosen Paper: Energy function **Mathematically**

A = Segmentation
binary vector

$$E(A) = \lambda \cdot R(A) + B(A)$$

Where **λ** is a weighting coefficient

Where **$R(A)$** is the regional term

$$R(A) = \sum_{p \in P} R_p(A_p)$$

Where **R_p** represents the penalty to assign to
a label **A_p**

e.g. could be represented as:

$$R_p(Obj) = -\ln \Pr(Ip | O)$$

$$R_p(Bkg) = -\ln \Pr(Ip | B)$$

Where **\Pr** is a histogram intensity distribution

Where **$B(A)$** is the boundary properties term

$$B(A) = \sum_{\{p, q\} \in N} B_{\{p, q\}} \cdot \delta(A_p, A_q)$$

Where **$B_{\{p, q\}}$** ≥ 0 :

Represents the similarity between pixels p and q.

The closer to 0, the less it is similar.

e.g. local intensity gradient, zero-crossing ...

And where **δ** function is:

$$\delta(A_p, A_q) = \begin{cases} 1 & A_p \neq A_q \\ 0 & \text{otherwise.} \end{cases}$$

Chosen Paper: Graph Initialization Mathematically

How is the **graph** initialized?

Graph Initialization

	<u>edge</u>	<u>weight (cost)</u>	<u>for</u>
neighbor links	$\{p, q\}$	$B_{\{p, q\}}$	$\{p, q\} \in N$
terminal links	$\{p, S\}$	$\lambda R_p(\text{"bkg"})$	$p \in P, p \notin O \cup B$
		K	$p \in O$
		0	$p \in B$
terminal links	$\{p, T\}$	$\lambda R_p(\text{"obj"})$	$p \in P, p \notin O \cup B$
		0	$p \in O$
		K	$p \in B$

selected by user

selected by user

And K is:

$$K = 1 + \max_{p \in P} \sum_{q: \{p, q\} \in N} B_{\{p, q\}}$$

What happens if we **assign more labels**?

Newly seeded pixels as object

<u>t-link</u>	<u>Initial cost</u>	<u>add</u>	<u>New cost</u>
$\{p, S\}$	$\lambda R_p(\text{"bkg"})$	$K + \lambda R_p(\text{"obj"})$	$K + C_p$
$\{p, T\}$	$\lambda R_p(\text{"obj"})$	$\lambda R_p(\text{"bkg"})$	C_p

Chosen Paper: Graph Initialization Mathematically

Once the graph is (re)initialized, to obtain the segmentation we can apply a low-order polynomial graph cut algorithm such as:

- Ford-Fulkerson style “augmenting paths”
 - Dinic algorithm $\rightarrow O(mn^2)$
 - Boykov-Kolmogorov version of Dinic’s $\rightarrow O(mn^2 |C|)$
- Goldberg-Tarjan style “push-relabel”


Side-note: In the paper, authors use Boykov-Kolmogorov algorithm because in practice is faster than other alternatives.

P1: Inpainting


Inpainting:

- Technique in which damaged, deteriorated or missing regions of images are reconstructed by interpolation of surrounding areas

before



after




before




after




original



mask



result



original



mask



result




Conclusions and Results

- Notable results for inpainting small regions
- Bigger region, bigger blur in inpainted region, (may occur because):
 - The content available for inpainting is partial
 - The # of neighbours pixels to look at is not enough (4-connectivity)
- The smooth transition of the Laplacian operator is not enough to “hide” a visible region in an optimal way

P2: Poisson editing

Poisson editing:

- Seamless insert an image into a destination, by shifting part of its environment (color) and keeping intact (or not) the original details (without “any” visually unappealing seams)



Source image B



Destination image A



Importing Gradients



Mixing Gradients

Source

Target

Result

Conclusions and Results

- Importing Gradients
 - Keeps details of the inserted image
 - Blurred background and visible borders
- Mixing Gradients
 - Keeps most relevant (striking) details of both images
 - Useful for objects with holes or transparent
- Choose the method based on the application





Source



Target



Result







P3: Segmentation:

Segmentation:

- Segmentation is the process of finding different **regions** or **segments** (set of pixels) that partitionate an image into **meaningful parts**.

Conclusions and Results

- The **Chan-Vese** algorithm is a suitable algorithm for a basic binary segmentation problem.
- Varying hyperparameters such as **length penalty**, **regularization term** will result in smoother segmentation or sharper boundaries



Discussion

Optimization for computer vision:

- Optimization is the **backbone** of most **algorithms** that solve **computer vision** tasks, therefore, is a **fundamental** field
- By itself, it can already output **acceptable solutions** to some computer vision tasks
- **No need** (or few) for **data** needed to operate, in **contrast** of **ML** or **DL** solutions
- **Easy** to explain or **understand**, in **contrast** to not-so-easy explainable **ML** or **DL** solutions
- As the task **difficulty** increases, generally, **more criteria(s)** needs to be defined, which can be a **challenging** task

Conclusions

Through this Module:

- Learnt the **optimization foundations** in the context of **computer vision**
- **Developed** and/or explained five different optimization algorithms for three different computer vision **tasks**
- Grasped an **understanding** of the **trade-off** between **explainability** and **performance**