Master in
Computer Vision
Barcelona

M3 – Machine Learning for Computer Vision

Project: Deep learning classification - **Final Presentation**

Group 7: Guillem Capellera, Johnny Nuñez and Anna Oliveras

Outline

- Week 1: Bag of Visual Words framework
- Week 2: Beyond BoVW → SVMs, Spatial Pyramids, Fisher Vectors

Handcrafted methods: Bag of Visual Words

- Week 3: From hand crafted to learnt features
- Week 4: Fine tuning of pre-trained CNNs → Densenet121
- Week 5: Training a CNN from scratch
  - Initial CNN
  - CNN Refinement
  - Residual Connections
  - Residual Network
  - Recap

Data driven methods: Deep Convolutional Networks

- Conclusions

Datasets

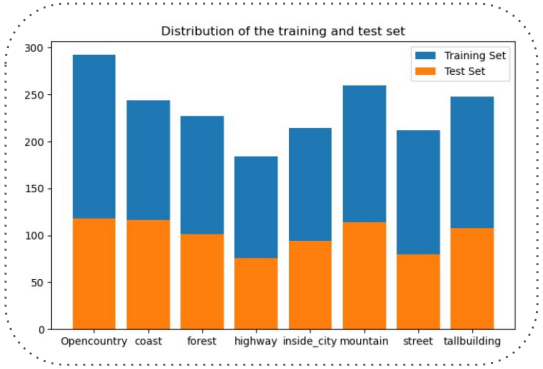

- We have 8 classes: coast, forest, highway, inside city, mountain, open country, street, tall buildings
- Big dataset : MIT_split → total of 2288 images
- Small dataset: MIT_small_train_1 → Train with only 50 images for each class !

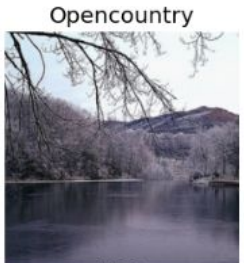

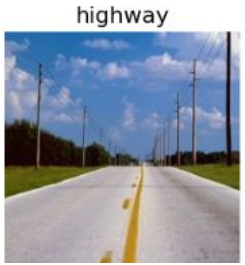

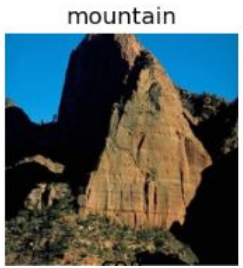

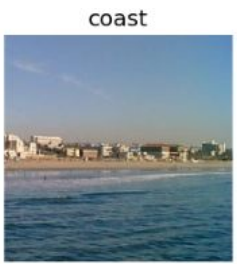

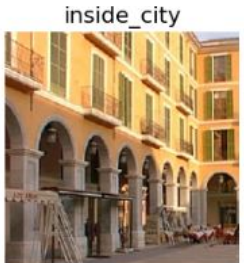

Week 1: Bag of Visual Words framework

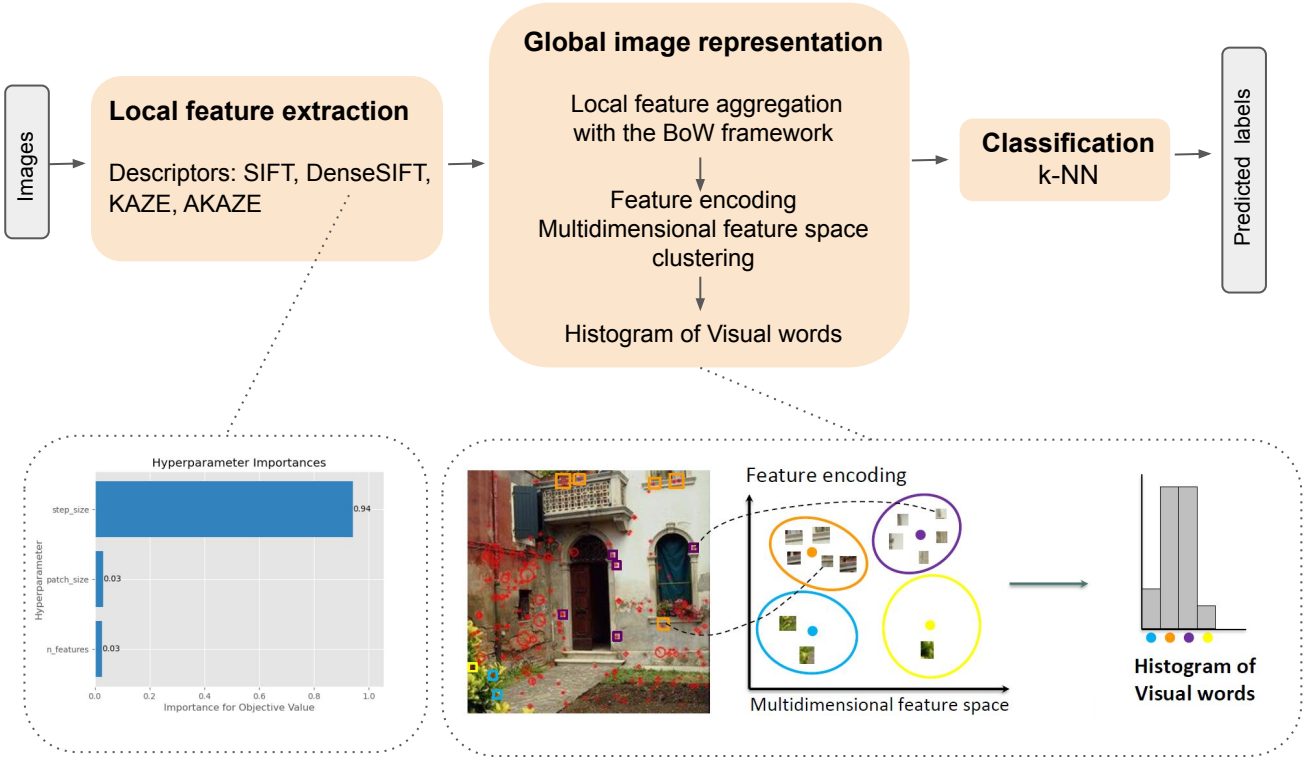

Grid search best parameters (Optuna)

Descriptor: DenseSIFT with
- n_features = 251
- patch_size = 3
- step size = 75

Clustering
- Codebook size k = 1024

Dimensionality reduction
- PCA with n_components = 46

Classifier
- n_neighbors = 18
- distance = euclidean

Accuracy = 0.95 **f1-score** = 0.8

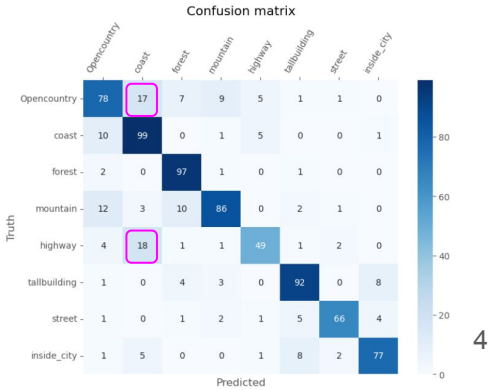

Week 2: Beyond BoVW

Spatial Pyramids

Multiple divisions:
1x1, 2x2, 4x4
vertical 1x2, 1x3, 1x4
horizontal 2x1, 3x1, 4x1

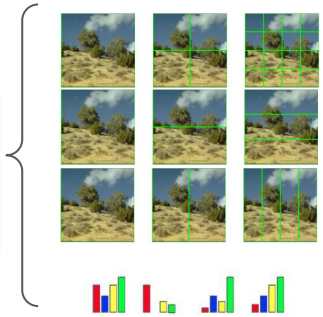

SVM and Kernels

- Linear
- RBF
- Histogram intersection

Feature normalization

- Power Norm
- L2 Norm

Fischer Vectors

Inlude higher order statistics: mean, covariance of local descriptors → GMM clustering

Accuracy = 0.96 (Without hyperparameter optimization)
F1 score = 0.83

Grid search best parameters (Optuna)→ We use Dense SIFT

- num features = 178
- **step size = 18**
- num cluster = 798
- num components = 69
- gamma = 0.00445
- C = 4.38
- **dim_reduc = PCA**
- **kernel = RBF**
- Normalization = power
- **Spatial_pyramid = vertical 1x4**

Accuracy = 0.96
F1 score = 0.86

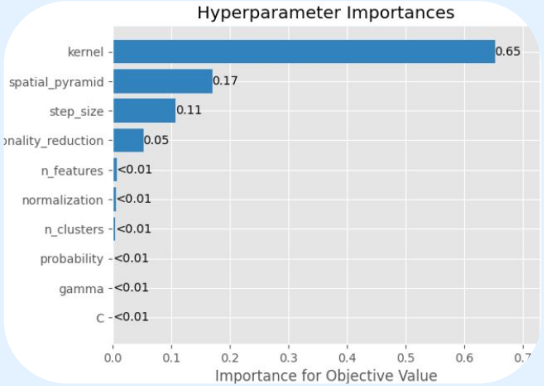

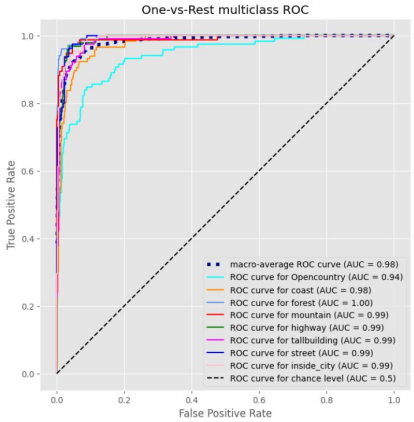

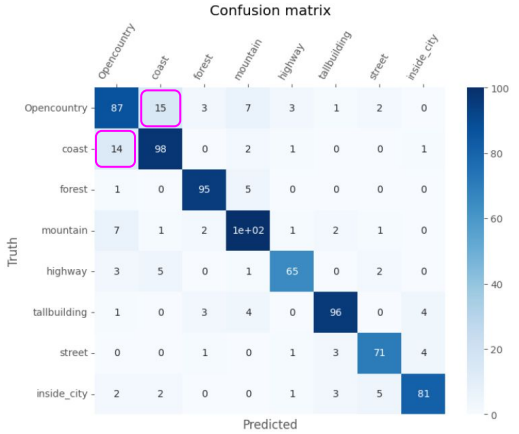

Week 3: From hand crafted to learnt features

Initial Model

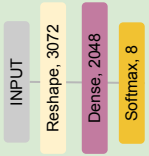

- 6M parameters
- Accuracy = **0.61**

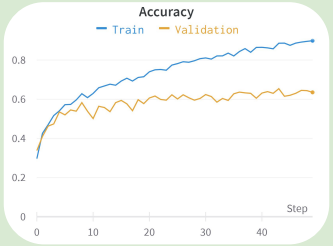

Best Model

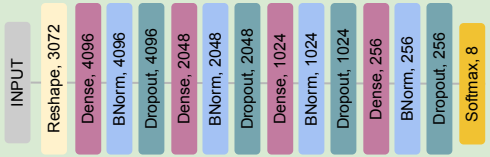

- 23M parameters
- Accuracy = **0.596**

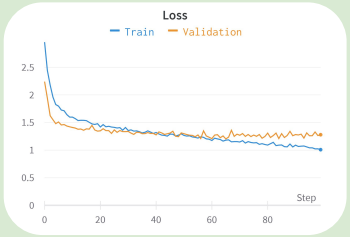

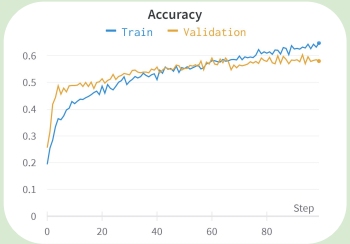

Deep features + SVM

Extracting the features after each dense layer output → similar results (best with output dense 4096)

- Accuracy = **0.41**

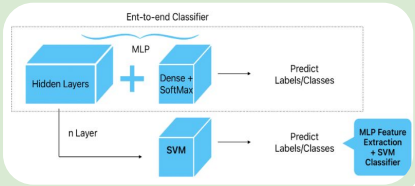

Patch based MLP

→ Best patch size: 32 x 32
→ Best aggregation: mean

- Accuracy = **0.77**

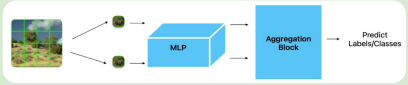

Patch based deep features + BoVW

→ Num clusters too low —> too much generalization of the features
→ Num clusters too big → too much specificity of the features
→ Best number of clusters (codebook size) = 256

- Accuracy = **0.72**

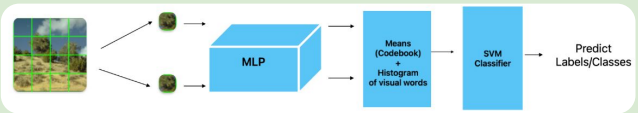

Week 4: Fine Tuning DenseNet121[1]

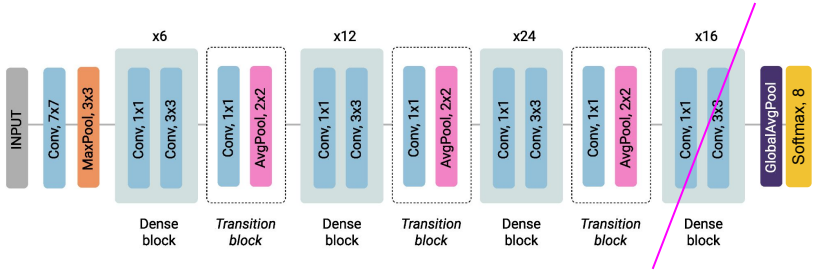

29% less parameters than the original model

Making the network smaller

MIT_split dataset

Model	Epochs	Num parameters	Validation accuracy
Original	300	7M	0.9542
Removing 1 DB	**300**	**5M**	**0.941**
Removing 2 DB	300	1.5M	0.825

Using MIT_small dataset

- **Overfitting**
- Accuracy drops to **0.845**

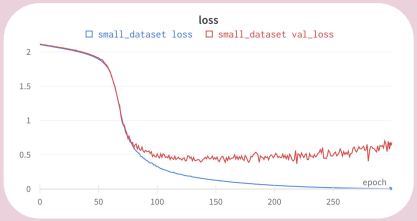

Data augmentation

- Horizontal Flip = True
- Zoom Range = 20%
- Accuracy increases to **0.895**
- Validation loss does not increase

Improve learning curve

- Early Stopper
- Reduce LR
- BatchNormalization and Dropout
  → Accuracy increases to **0.915**

Hyperparameter Optimization

→ Accuracy increases to **0.9518**

Optimizer = Adamax
LR = 0.0001
Dropout = 0.5

Wight decay = 0.3
Batch_ size = 10
BatchNormalization

[1] Huang, G., Liu, Z., Weinberger, K. Q., & van der Maaten, L. (2017). Densely connected convolutional networks. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (pp. 4700-4708). https://arxiv.org/abs/1608.06993

Building a CNN from scratch

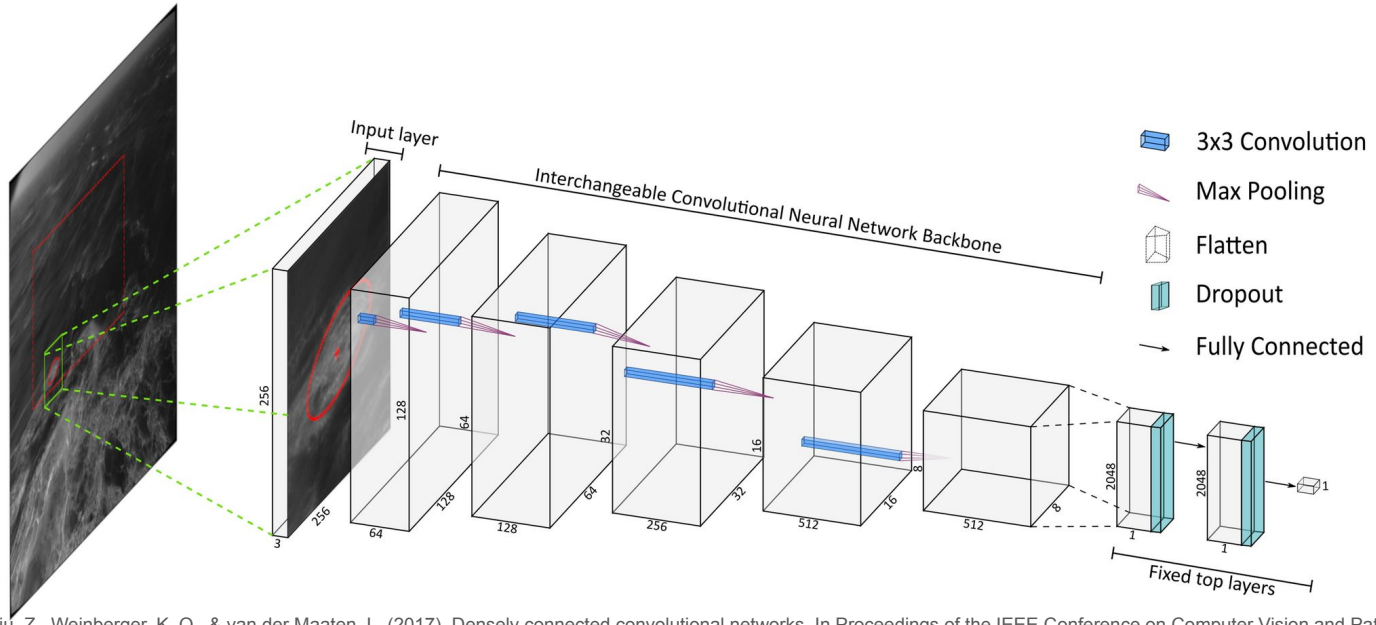

[1] Huang, G., Liu, Z., Weinberger, K. Q., & van der Maaten, L. (2017). Densely connected convolutional networks. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (pp. 4700-4708). https://arxiv.org/abs/1608.06993

Manual search of the CNN network

Params	Accuracy
1.5M	0.75

- CNN
  - 4 blocks → Conv2D + Conv2D + MaxPool + BatchNorm
- Classifier
  - Global Average Pooling
  - 3 blocks → Dense + Dropout

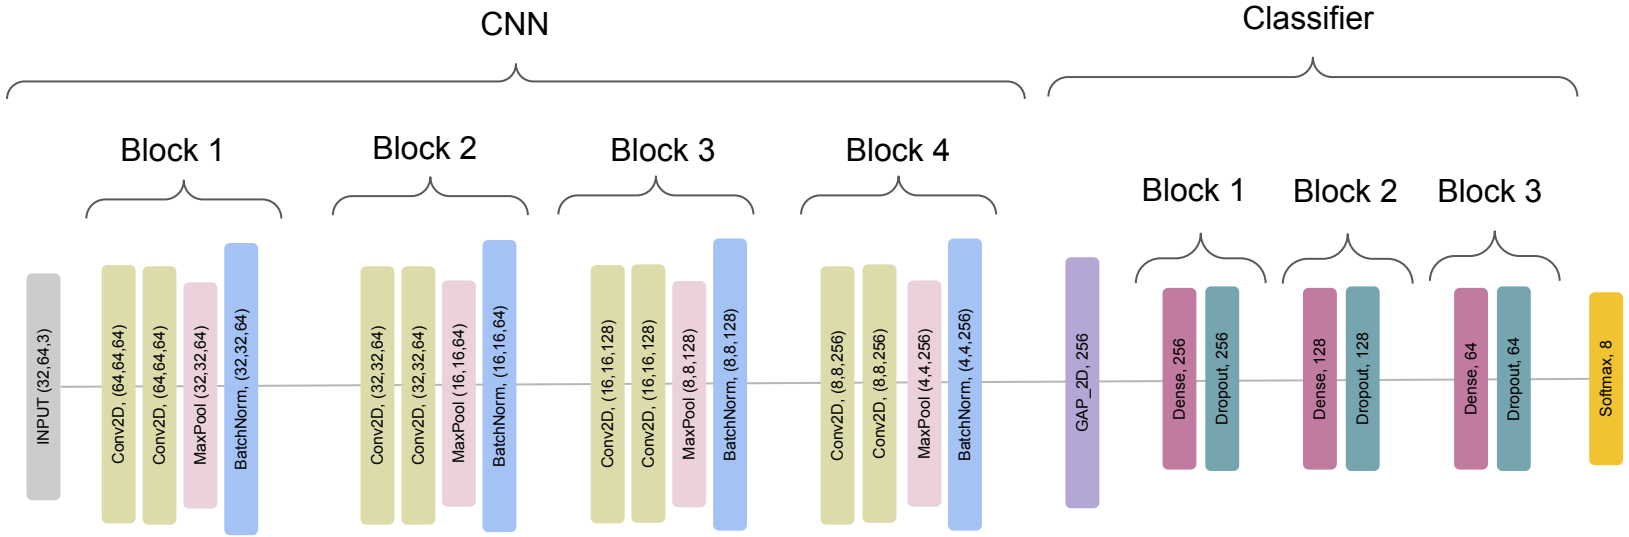

Network refinement

Params	Accuracy
1.3M ↓	0.75

- Number of CNN blocks → 2, 3, **4**, 5
- For each CNN block → Add a second layer? **True**, False
- **Dimensionality first filter** → **32**, 64
- Number of dense layers → 1, 2, **3**, 4

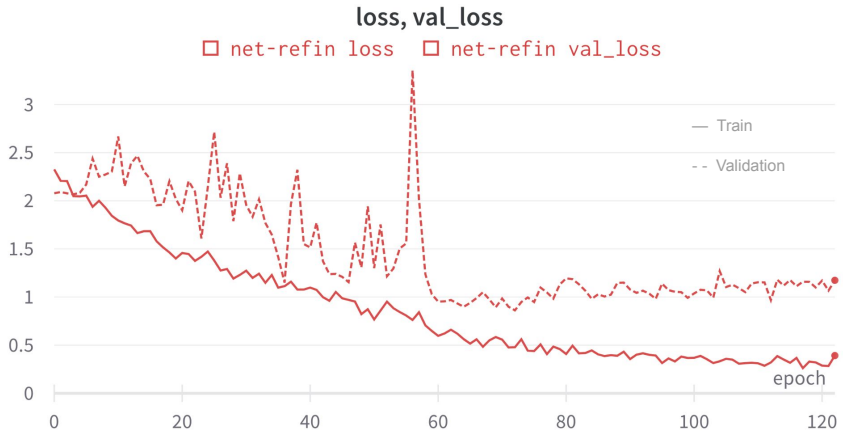

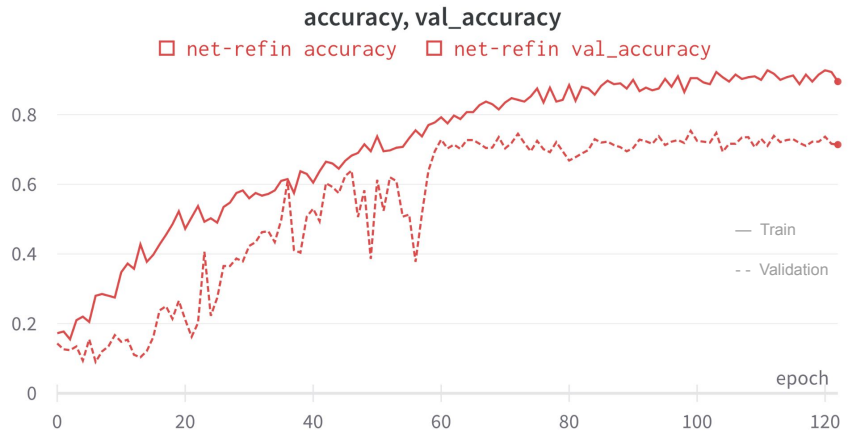

Initializer & activation function

Params	Accuracy
1.3M	0.81 ↑

- Initializer → GlorotUniform, GlorotNormal, HeUniform, HeNormal
- Activation → Relu, LeakyRelu, Mish
- Kernel size → 3x3 or 5x5
- Kernel regularizer → True, False

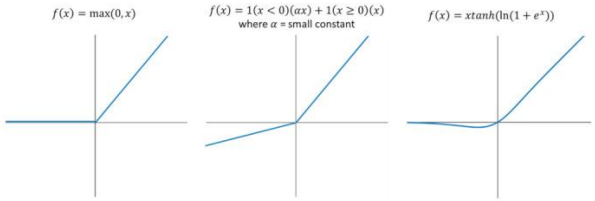

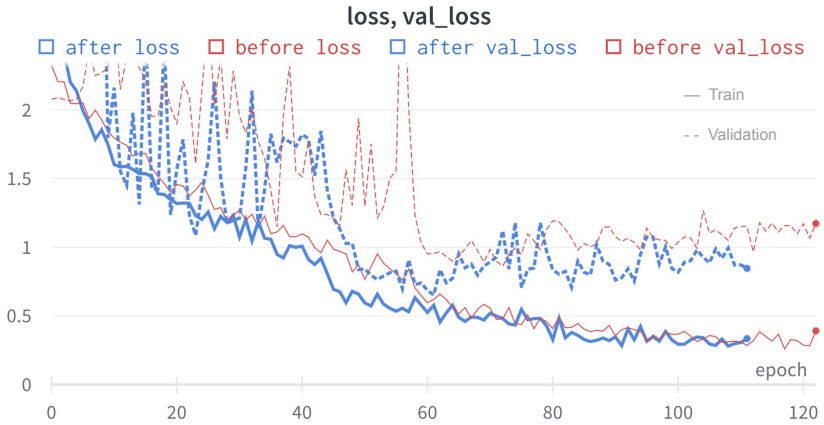

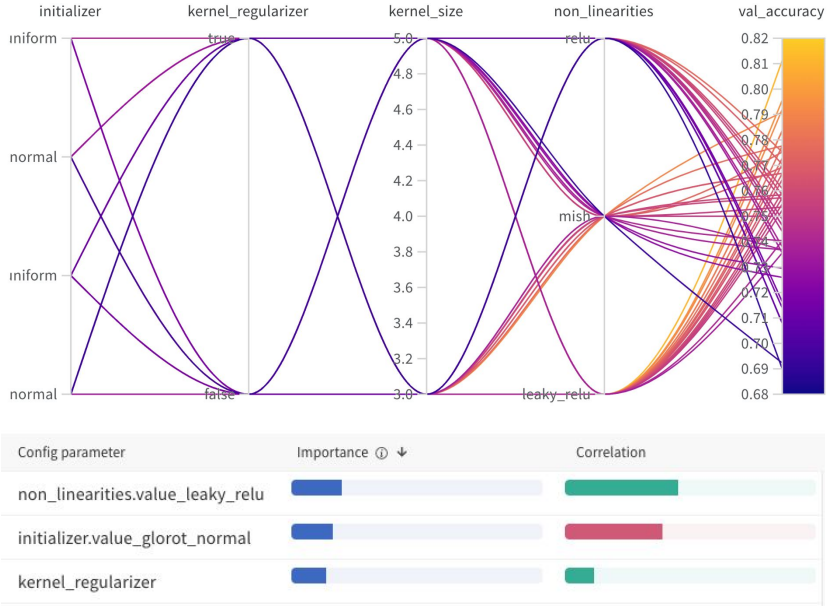

Config parameter	Importance ⓘ ↓	Correlation
non_linearities.value_leaky_relu		
initializer.value_glorot_normal		
kernel_regularizer		

Initializer & activation function

Params	Accuracy
1.3M	0.81 ↑

- Initializer → GlorotUniform, GlorotNormal, HeUniform, HeNormal
- Activation → Relu, LeakyRelu, Mish
- Kernel size → 3x3 or 5x5
- Kernel regularizer → True, False
-

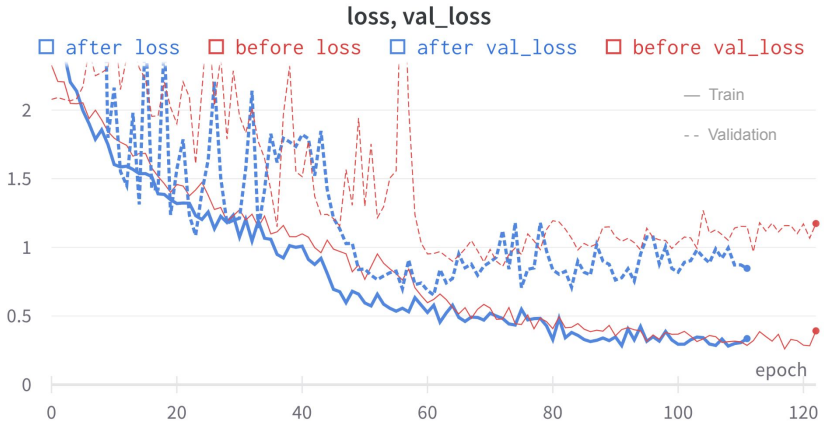

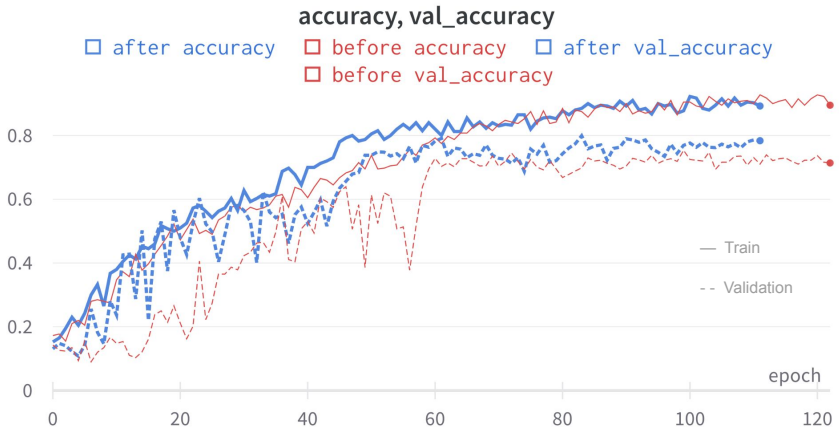

Hyperparameter optimization

Params	Accuracy
1.3M	0.85 ↑

Sweep	Values	Best	Sensitivity
LR	[0.01, 0.005, 0.001]	0.001	High
Batch Size	[8, 16, 32, 64]	32	Medium
Optimizer	[Adam, SGD]	Adam	Low
Dropout values	[0.3 , 0.4, 0.6]	0.4	Medium
Horizontal Flip	[True, False]	True	High
Rotation	[0, 15]	0	Medium
Width Shift	[0, 0.1]	0.1	Medium
Height Shift	[0, 0.1]	0.1	Medium
Shear Range	[0, 0.1]	0	Medium
Zoom Range	[0, 0.1]	0	High

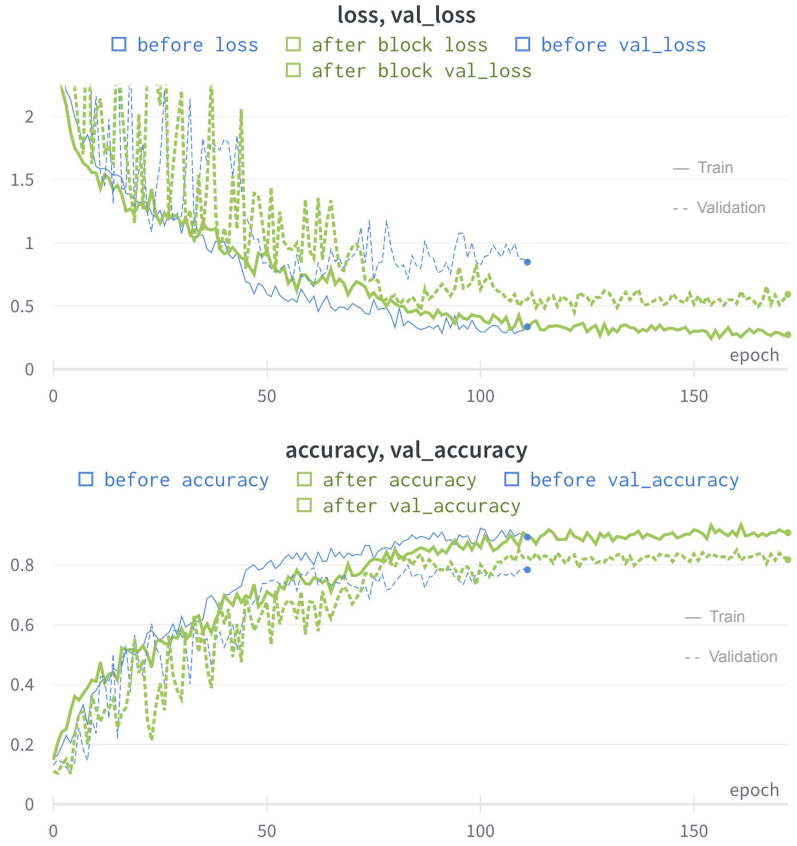

Parameter refinement

Params	Accuracy
1.3M	0.85

- Reduce the number of **filters** at each CNN block

Params	Accuracy
380k ↓ ↓	0.80 ↓

CNN blocks	Before	After
1st block	32, 32	32, 32
2nd block	64, 64	**32, 32**
3th block	128, 128	**64, 64**
4th block	256, 256	**128, 128**

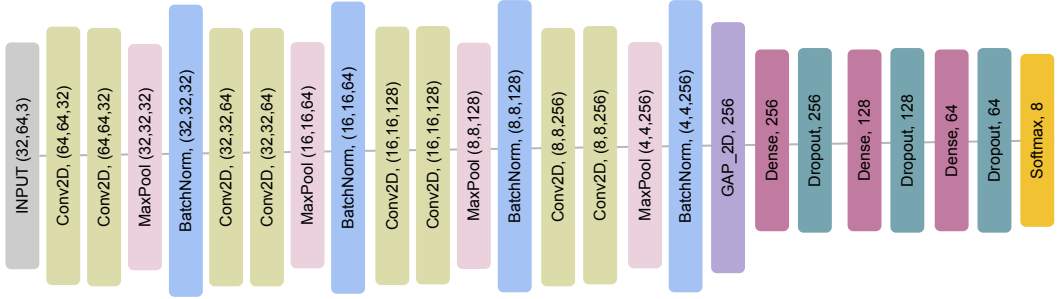

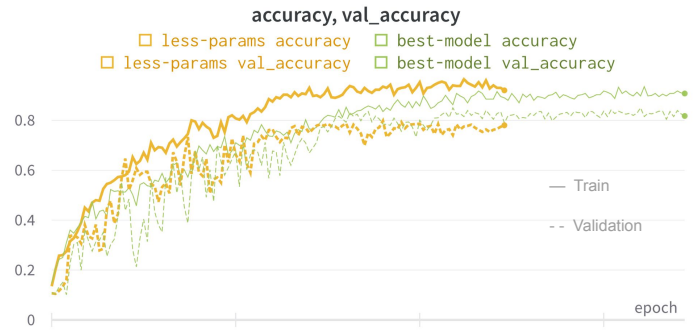

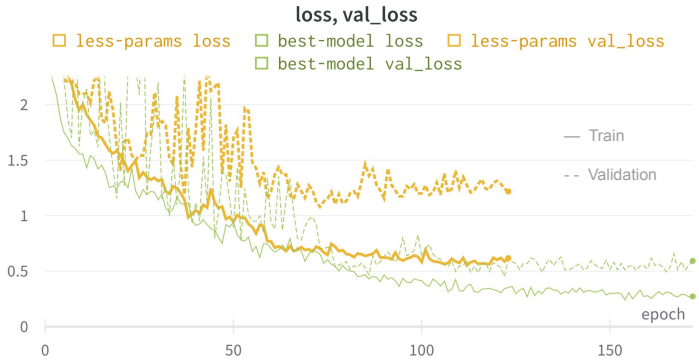

Residual Connections

- Residual Connections between blocks
- Reduce dimensionality of the filters

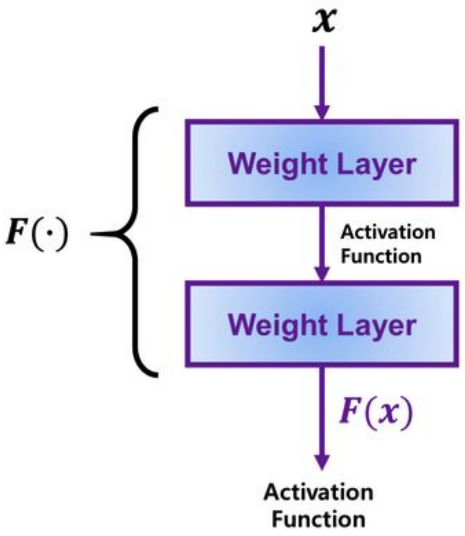

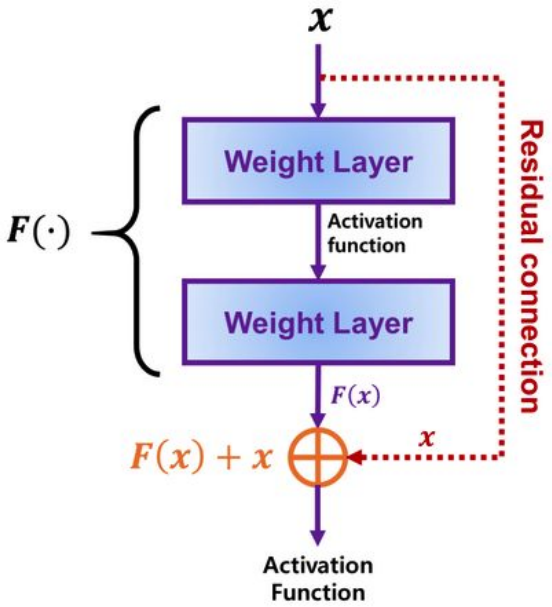

Benefits

- Converge more easily
- More Stability Training
- Better Generalization
- Easy implementation
- Improve accuracy

Residual network search

Params	Accuracy
68k	0.82

- CNN blocks → 2, 3, **4**, 5
- Num filters → 16, **32**, 64, 128, 256
- Dense layers → 0, **1**, 2, 3

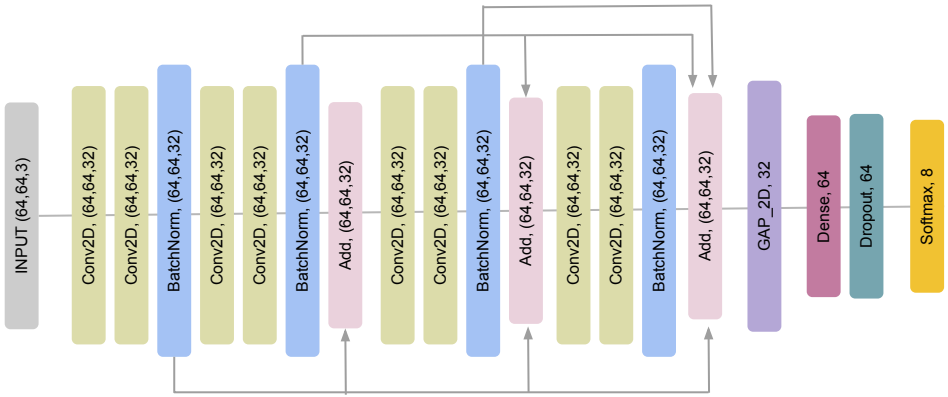

68k!!!

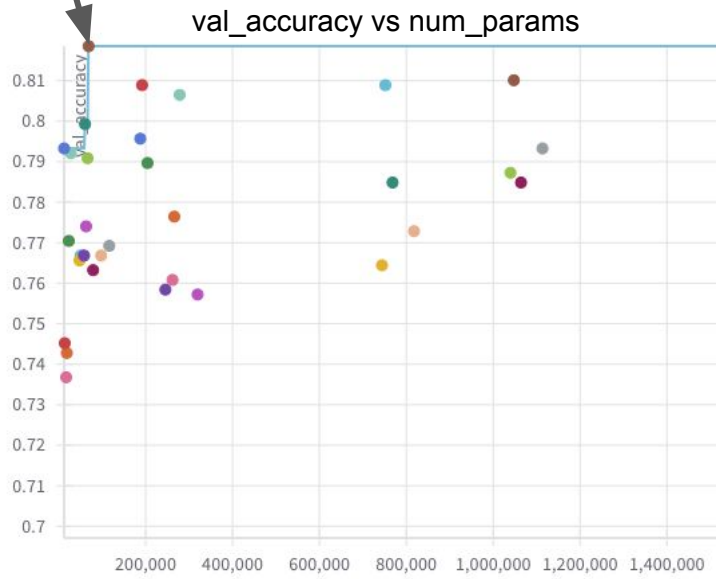

Residual network search

Params	Accuracy
68k	0.82

- CNN blocks → 2, 3, **4**, 5
- Num filters → 16, **32**, 64, 128, 256
- Dense layers → 0, **1**, 2, 3

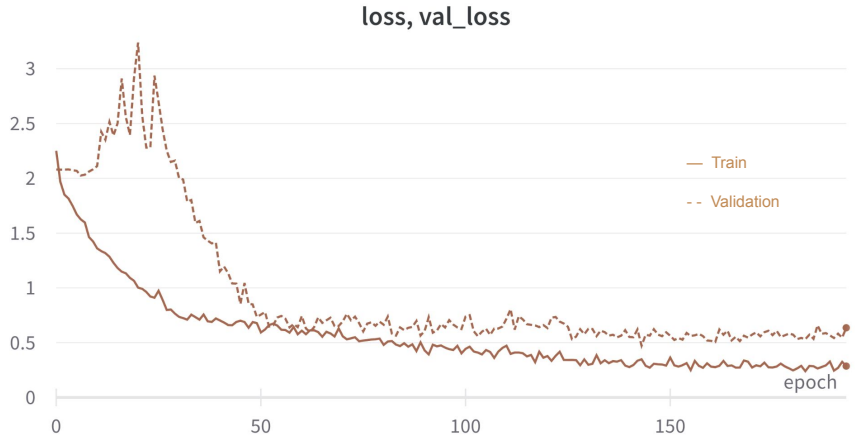

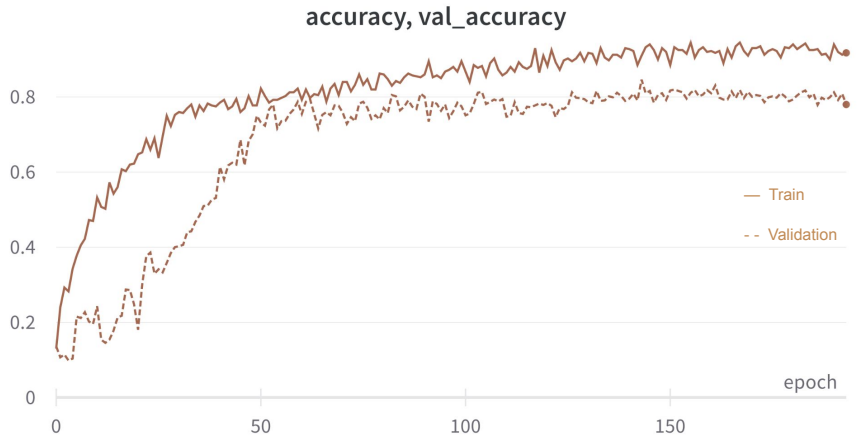

Data augmentation refinement

Params	Accuracy
68k	0.85 ↑

- Horizontal Flip → True
- Rotation [0, 15] → **0**
- Width Shift [0, 0.1] → **0.1**
- Height Shift [0, 0.1] → **0**
- Shear Range [0, 0.1] → 0
- Zoom Range [0, 0.1] → 0

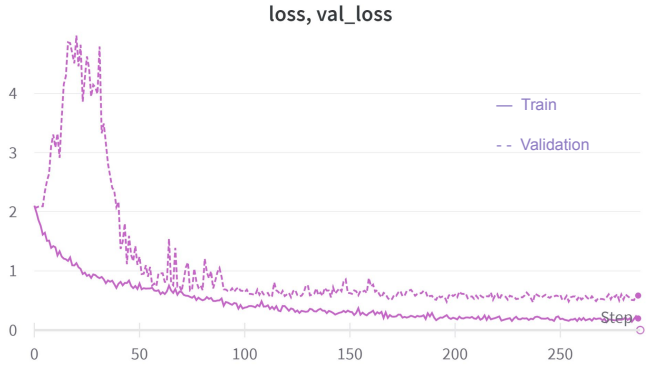

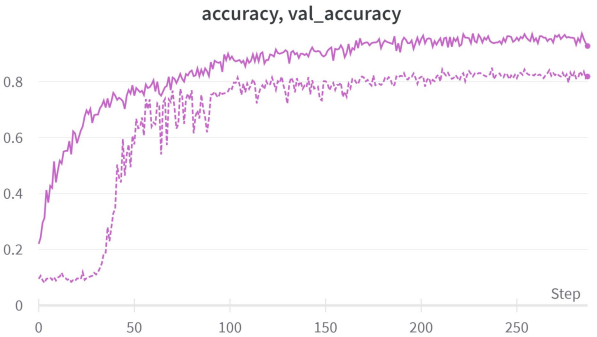

Recap week 5

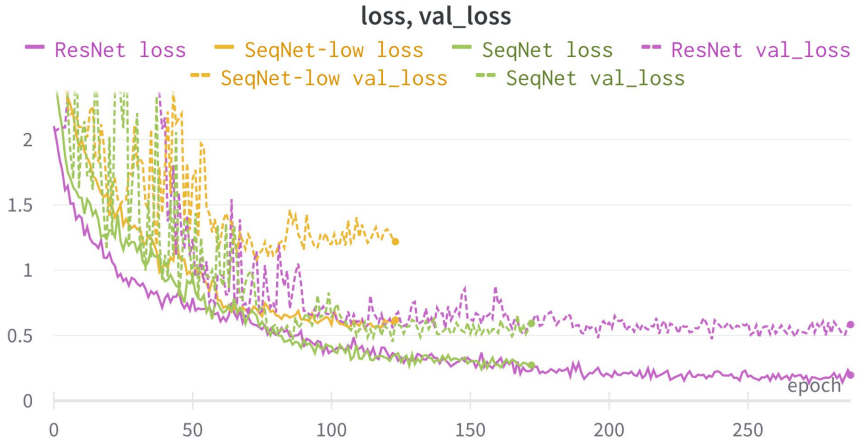

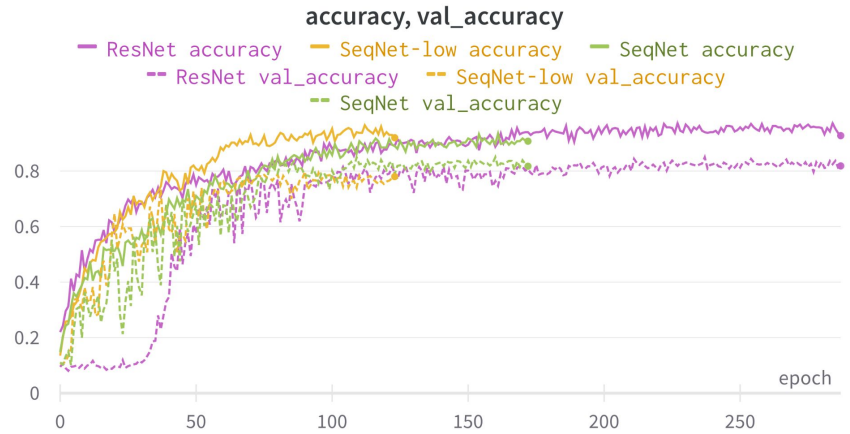

Params	Accuracy
1.3M	0.85

Params	Accuracy
380k	0.80

Params	Accuracy
68k	0.85

Recap week 5

- Good results with BOVW and spatial pyramids (**0.95** accuracy)
- MLP → too simple for our problem in all the cases (**0.77** accuracy)
- Reduced data → challenging and requires the use of **data augmentation** and other techniques to achieve the desired results
- Fine-tuning the DenseNet121 → best results (**0.96** accuracy) and we were able to reduce **30%** of the network parameters to **5M**
- Building a network from scratch is time consuming and difficult to optimize
- Residual connections → help the network to converge easily and generalize.

  → **0.85** accuracy with **68k** parameters
- Our results are limited due to the lack of data