



# Master in Computer Vision *Barcelona*

Image Classification

Ramon Baldrich

# Artificial Intelligence in Computer Vision

Machine learning for image classification:

Data driven methods: Deep Convolutional Networks: 3 sessions

From hand-crafted to learnt features

Fine tuning of pre-trained CNNs

Training a CNN from scratch



# Xception 0.92

configuration	train accuracy	test accuracy
frozen backbone	.98	.893
exit flow	1	.924
block 8	1	.926
block 7	1	.924
block 6	.97	.931
block 5	1	.926
block 4	1	.920
block 3	1	.917
block 2	1	.931
block 1	1	.930
train whole model	1	.909




configuration	train accuracy	test accuracy
dropout (0.1)	.993	0.908
BatchNorm	.987	.907
L1 regularizer	.994	.931
L1 + dropout	.998	.929
L1 + BatchNorm	.989	.923
BatchNorm + dropout	.986	.869
L1 + BatchNorm + Dropout	.989	.887


	accuracy	loss
training (baseline)	0.901	0.448
testing (baseline)	0.92	0.431
training (optuna)	0.98	0.07
testing (optuna)	0.979	0.09

We saw how in general it is better to finetune some of the feature extractor's layers on our dataset, instead of keeping it entirely frozen. Nevertheless, it is not advisable to totally unfreeze it.

# VGG16 0.89



# ResNet50 ??

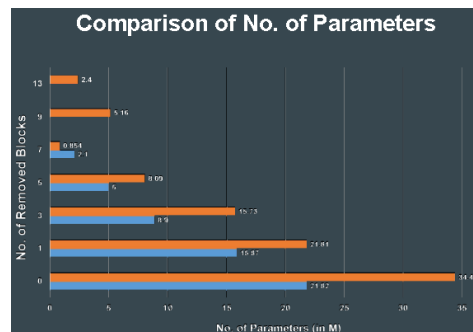




Better use train/val loss

Test accuracy after training		Added Dropout after every conv	
Dropped Blocks		Accuracy	Loss
		0.5717	1.2997
		0.7954	0.6399
		0.8899	0.3545
		0.9380	0.2162
		0.9318	0.2031



# Inception v3 0.85-ish

Used Augmentation	Values	Accuracy (V3)	Accuracy (V4)
Horizontal flip	True	0.843	0.885
Vertical flip	True	0.822	0.885
Rotation	30	0.854	0.823
Rotation	10	0.875	<b>0.906</b>
Zoom	0.2 - 52	0.375	0.865
Feature-wise center	True	0.864	0.875
Feature-wise STD norm	True	<b>0.885</b>	0.896
Shear range	0.2	<b>0.885</b>	0.875
Feature-wise center + Feature-wise STD norm	True, True	0.875	0.885
Horizontal flip + Vertical flip	True, True	0.844	0.843





# InceptionResnetV2 0.92

All frozen until last one



Val. acc.: 0.92


No one frozen



Val. acc.: 0.88


Removing blocks method 1:

Removed blocks	Blocks left	Partition of remove (per group)	Parameters	Validation Accuracy
0	40		54 M	87,6
4	36	1 2 1	50 M	89,7
8	32	2 4 2	45 M	89,5
16	24	4 8 4	36 M	89,1
32	8	8 16 8	19 M	88,9
36	4	9 18 9	14 M	89,1



Removing blocks method 2:


Removed group	Parameters	Validation Accuracy
Block8-s	32785384	0.92
Block17-s	5478408	0.91
<b>Block35-s</b>	<b>567896</b>	<b>0.90</b>



We achieved a validation accuracy of 89% with a network having 500k parameters (108 times less parameters than the original network) using only 400 training images

# MobileNet 0.91

Same accuracy removing 2 blocks



~88% Validation Accuracy on small dataset

Focus on the message

[https://wandb.ai/mcv-m3-g6/task\\_1\\_sweep/reports/Slides-M3-Week-4-Team-6--VmlldzozNDMONTY1?accessToken=poxugew27pus1iukp4o5e39I53d8Irf0by30c6tugqfirps0dnin4rwq7gz2tp6w](https://wandb.ai/mcv-m3-g6/task_1_sweep/reports/Slides-M3-Week-4-Team-6--VmlldzozNDMONTY1?accessToken=poxugew27pus1iukp4o5e39I53d8Irf0by30c6tugqfirps0dnin4rwq7gz2tp6w)





# DenseNet121 0.917

Model	Epochs	Num parameters	Validation accuracy
Original	50	7M	0.9318
Original	300	7M	0.9542
Removing 1 DB (x16)	300	5M	0.941
Removing 2 DB (x24.x16)	300	1.5M	0.825
Removing 3 DB (x24.x16.x12)	300	380.000	0.52
Removing 1 DB + adding dense [1024]	300	5M	0.9393
Removing 2 DB + adding dense [1024]	300	1.5M	0.832
Removing 3 DB + adding dense [1024]	300	520.000	0.601

Dataset	Validation accuracy
MIT	0.941
MIT_small_1	0.845
MIT_small_1 data augmentation	0.895


5%!!



Dropout worth the price

NasNetMobile 0.913


# EfficientNetB0 0.886



In order to set a new model we removed blocks 4, 5, 6 and 7


0.855 on small dataset

8 Batch		16 Batch		32 Batch		64 Batch	
Train acc	Val acc	Train acc	Val acc	Train acc	Val acc	Train acc	Val acc
96.24%	87.00%	94.50%	87.88%	98.09%	85.37%	91.07%	84.89%



# ConvNextTiny 0.934

## Overfits?




Training set	Accuracy
MIT_small_1	0.9165
MIT_small_2	0.9318
MIT_small_3	0.9335
<b>MIT_small_4</b>	<b>0.9419</b>
MIT_training(transfer learning)	0.9343

## Data aug.

After scaling down our search space the following data augmentation parameter values resulted in models with high accuracies:

- Height shift: <0.06
- Horizontal flip: True
- Rotation: <10 deg
- Shear shift: <0.28
- Vertical flip: False
- Width shift: <0.13
- Zoom: <0.1

D.A depends on the architecture









Group	grade
1	10
2	6
3	9
4	8
5	10
6	8
7	10
8	9
9	9
10	9

# EXTRA: Interpretability - Class Activation Maps

Correct highway

Correct street

Correct tallbuildings






# EXTRA: Interpretability - Errors analysis


(Ground truth => **Network prediction**)

to the appearance of key features  
of other classes in the image. This  
is very clear with opencountry.


Opencountry => **Mountain**




Tallbuilding => **Highway**




Coast => **Tallbuilding**



Forest => **Mountain**



Opencountry => **Coast**



Opencountry => **Inside\_city**

