



# Master in Computer Vision *Barcelona*

**Module 3:** Machine Learning for Computer Vision


**Project:** Deep learning classification

**Lecturer:** Ramon Baldrich

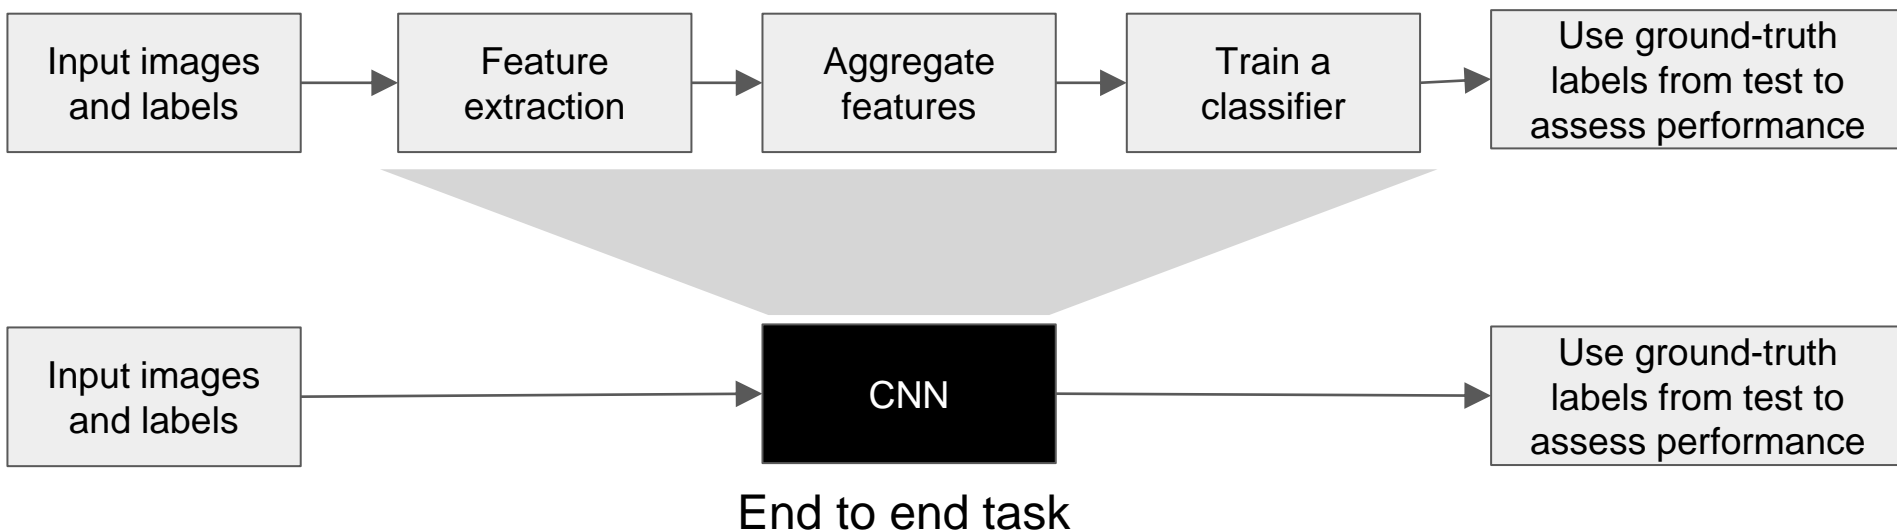


## Module Goal

The aim of this module is to learn the techniques for category classification: handcrafted and learned.



# Pipeline of the project W5 and W6



Machine learning for image classification:

Data driven methods: Deep Convolutional Networks: 3 sessions

From hand-crafted to learnt features

Fine tuning of pre-trained CNNs

Training a CNN from scratch

# Keras: first example

## # create model

```
model = Sequential()
model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
model.add(Dense(8, init='uniform', activation='relu'))

inputs = Input(shape=None))
x = Dense(12, init='uniform', activation='relu', name='fc1')(x)
x = Dense(8, init='uniform', activation='sigmoid', name='predictions')(x)
model = Model(inputs, x, name='example')
```

W3-5

## # Compile model

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

## # Fit the model

```
model.fit(X, Y, nb_epoch=150, batch_size=10)
```

## # evaluate the model

```
scores = model.evaluate(X, Y)
print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

W3-4

## # predict with the model

```
features = model.predict(X)
```

W3-4

# Understanding CNN topology: filtering

A guide to convolution arithmetic for deep learning

Vincent Dumoulin<sup>1\*</sup> and Francesco Visin<sup>2\*†</sup>


<sup>\*</sup>MILA, Université de Montréal

<sup>†</sup>AIRLab, Politecnico di Milano

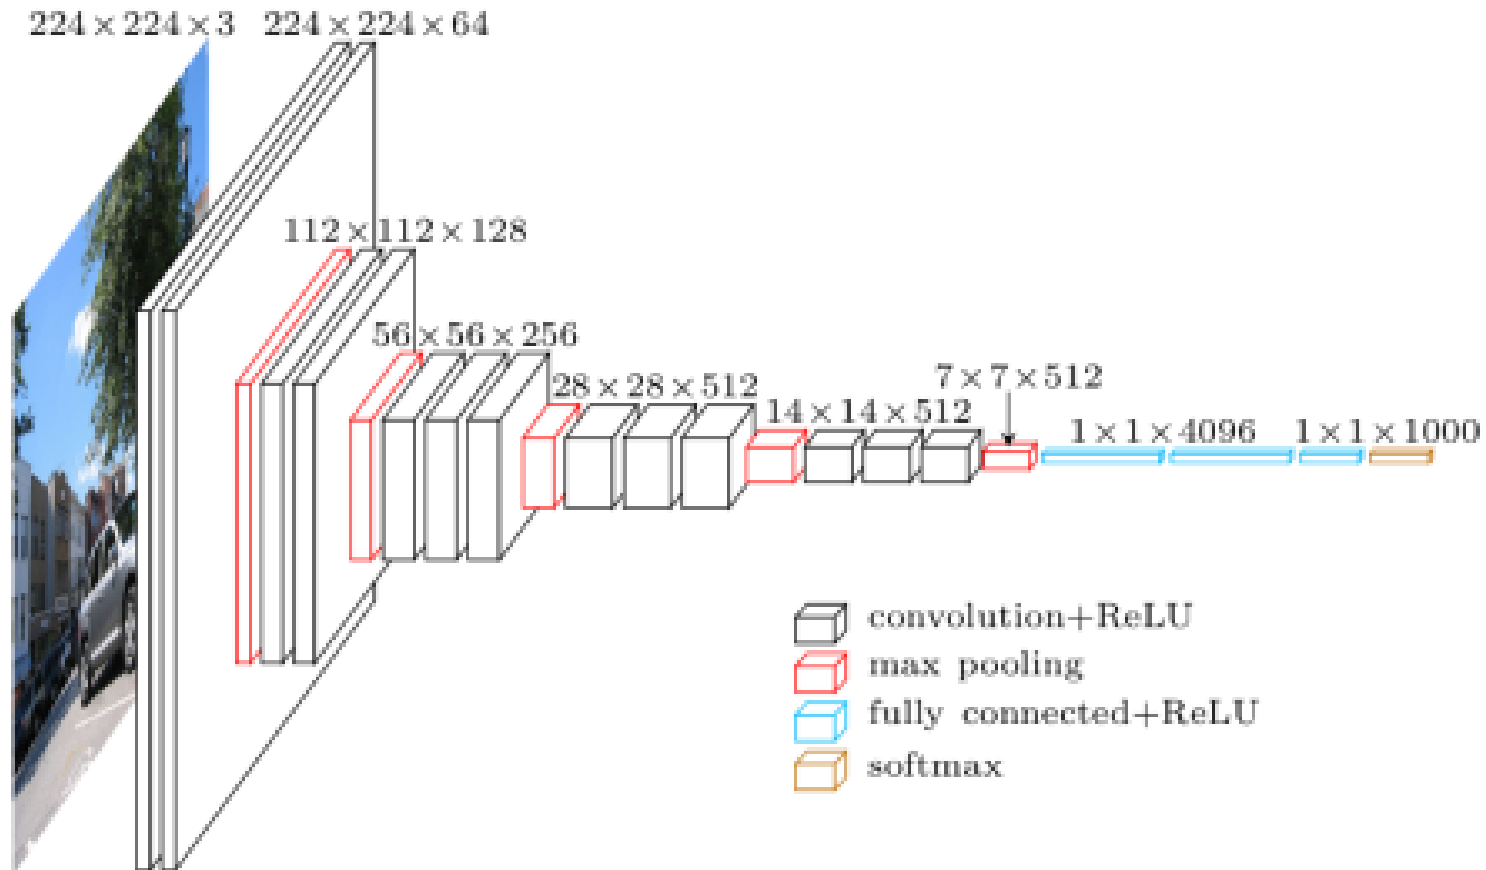
March 24, 2016

<https://arxiv.org/pdf/1603.07285v1.pdf>

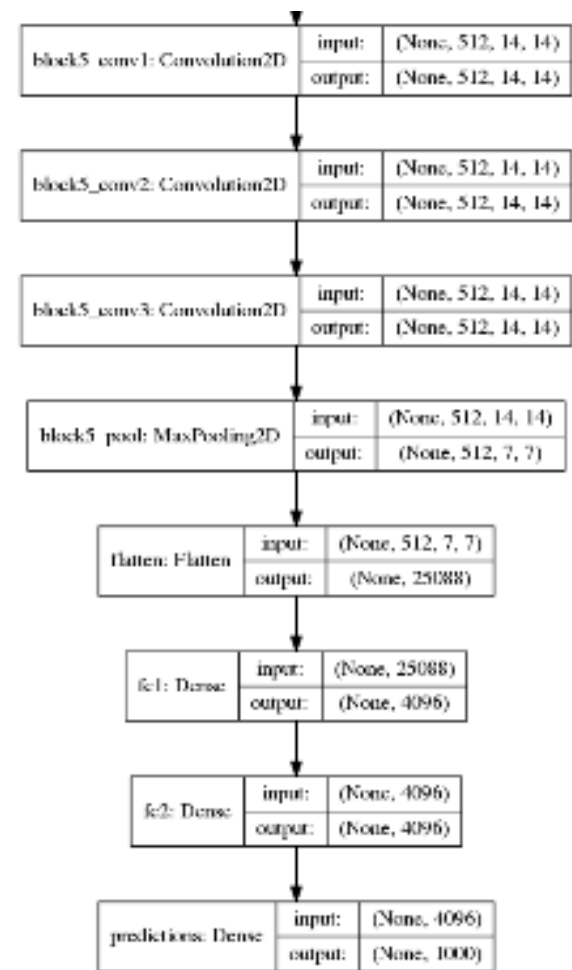
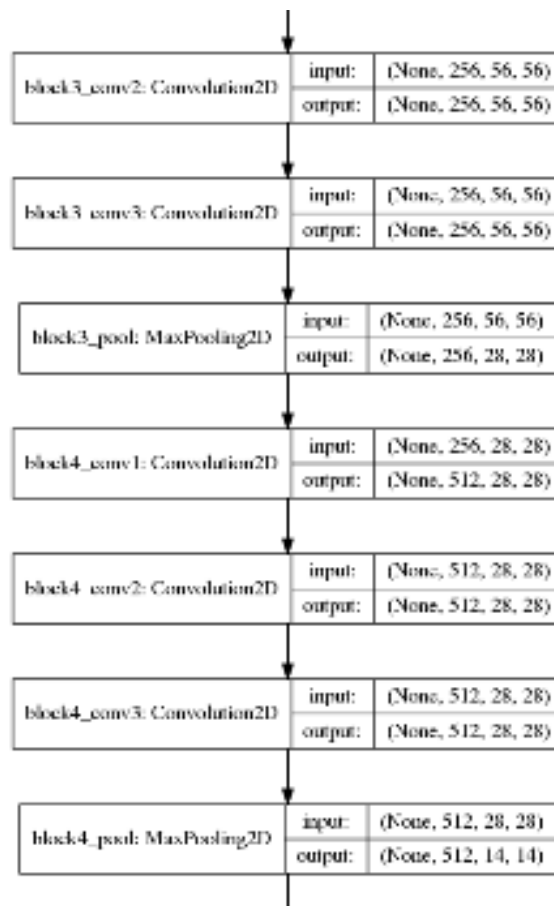
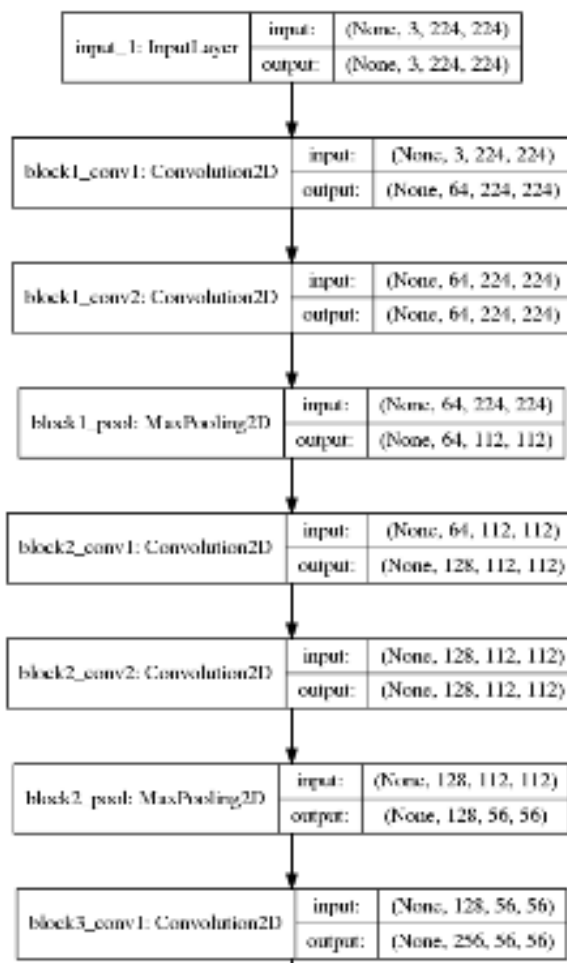
Input:  $5 \times 5 \times 2$   
Filter:  $3 \times 3 \times 3 \rightarrow 3 \times 3 \times 2 \times 3$   
Output:  $3 \times 3 \times 3$  or  $5 \times 5 \times 3$

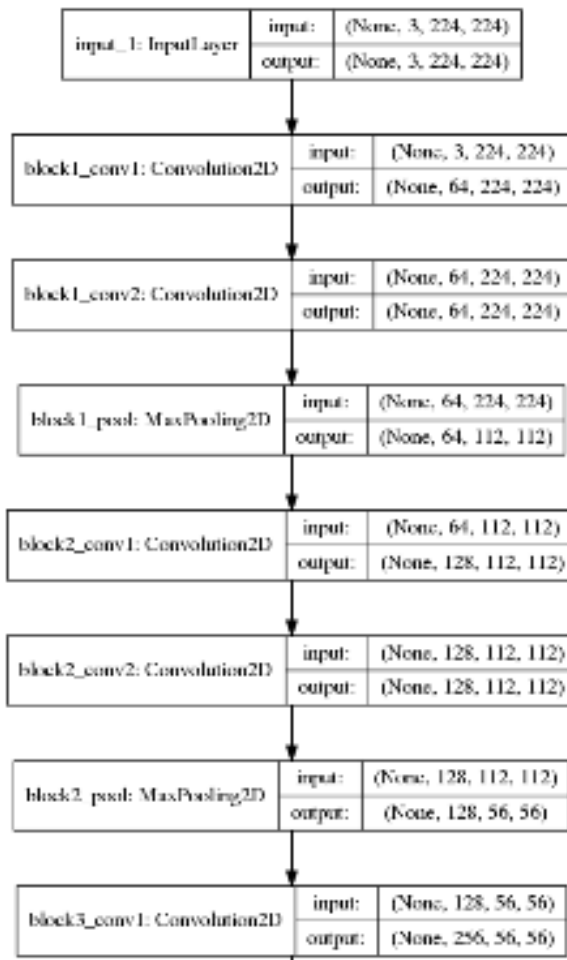


# Very deep convolutional networks for large-scale image recognition



Credit Davi Frossard





```
img_input = Input(shape=(3,224,224))
```

```
x = Convolution2D(64, 3, 3, activation='relu',  
border_mode='same', name='block1_conv1')(img_input)
```

```
x = Convolution2D(64, 3, 3, activation='relu',  
border_mode='same', name='block1_conv2')(x)
```

```
x = MaxPooling2D((2, 2), strides=(2, 2),  
name='block1_pool')(x)
```

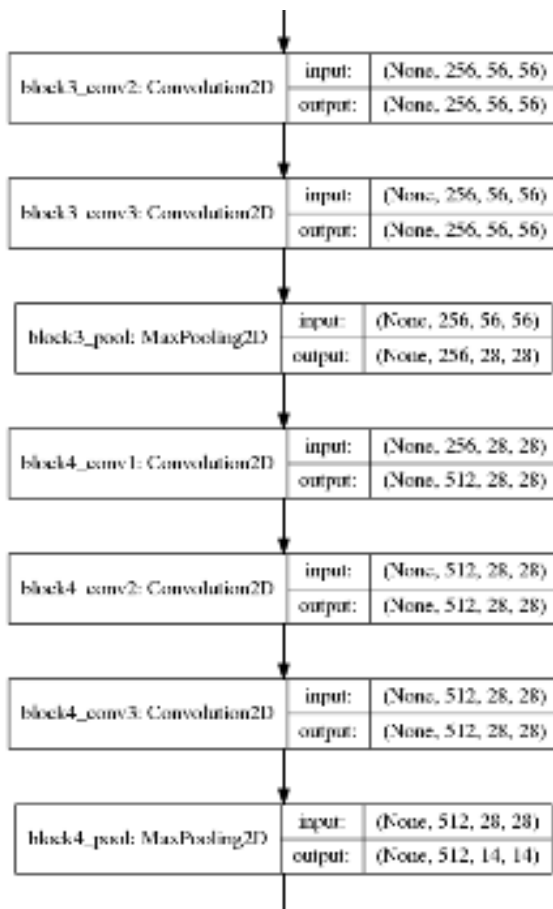
```
x = Convolution2D(128, 3, 3, activation='relu',  
border_mode='same', name='block2_conv1')(x)
```

```
x = Convolution2D(128, 3, 3, activation='relu',  
border_mode='same', name='block2_conv2')(x)
```

```
x = MaxPooling2D((2, 2), strides=(2, 2),  
name='block2_pool')(x)
```

```
x = Convolution2D(256, 3, 3, activation='relu',  
border_mode='same', name='block3_conv1')(x)
```





```
x = Convolution2D(256, 3, 3, activation='relu',
border_mode='same', name='block3_conv2')(x)
```

```
x = Convolution2D(256, 3, 3, activation='relu',
border_mode='same', name='block3_conv3')(x)
```

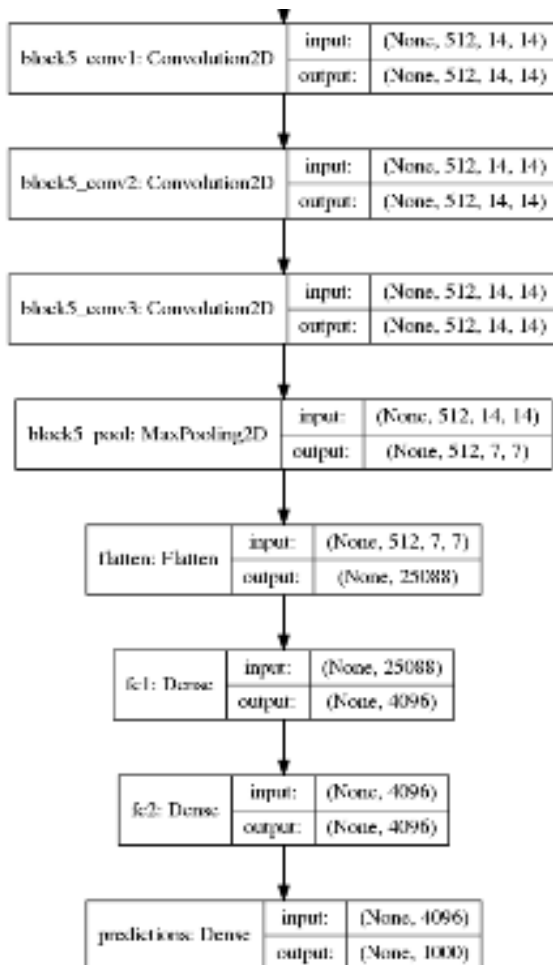
```
x = MaxPooling2D((2, 2), strides=(2, 2),
name='block3_pool')(x)
```

```
x = Convolution2D(512, 3, 3, activation='relu',
border_mode='same', name='block4_conv1')(x)
```

```
x = Convolution2D(512, 3, 3, activation='relu',
border_mode='same', name='block4_conv2')(x)
```

```
x = Convolution2D(512, 3, 3, activation='relu',
border_mode='same', name='block4_conv3')(x)
```

```
x = MaxPooling2D((2, 2), strides=(2, 2),
name='block4_pool')(x)
```



`x = Convolution2D(512, 3, 3, activation='relu', border_mode='same', name='block5_conv1')(x)`

`x = Convolution2D(512, 3, 3, activation='relu', border_mode='same', name='block5_conv2')(x)`

`x = Convolution2D(512, 3, 3, activation='relu', border_mode='same', name='block5_conv3')(x)`

`x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)`

`x = Flatten(name='flatten')(x)`

`x = Dense(4096, activation='relu', name='fc1')(x)`


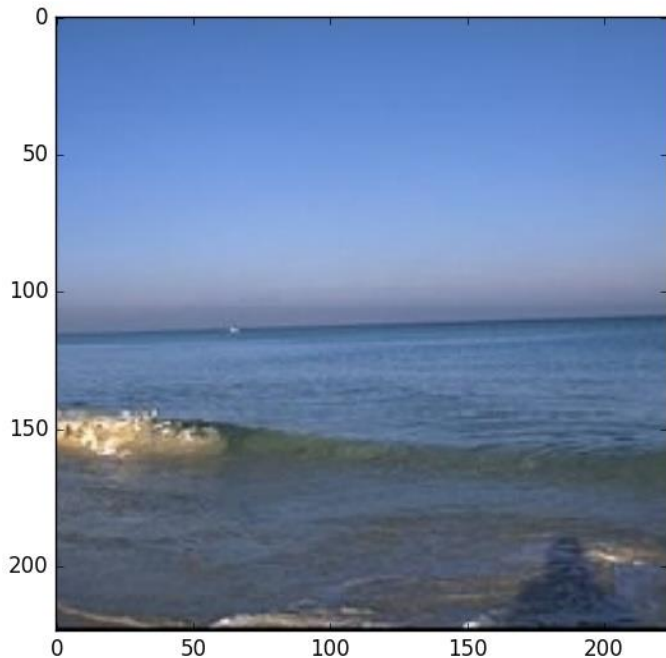
`x = Dense(4096, activation='relu', name='fc2')(x)`

`x = Dense(1000, activation='softmax', name='predictions')(x)`

# Extract features maps

```
img_path = '/data/MIT/test/coast/art1130.jpg'  
img = image.load_img(img_path, target_size=(224, 224))  
x = image.img_to_array(img)  
x = np.expand_dims(x, axis=0)  
x = preprocess_input(x)
```

```
base_model = VGG16(weights='imagenet')  
model = Model(inputs=base_model.input, outputs=base_model.get_layer('block1_conv1').output)  
features = model.predict(x)
```



For visualizing purposes:  
How to get rid of 3rd dimensión?

# Week 5: Fine tune end to end classification

1

## Return of the Devil in the Details: Delving Deep into Convolutional Nets

Ken Chatfield, Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman  
Visual Geometry Group, Department of Engineering Science, University of Oxford  
{ken,karen,vedaldi,az}@robots.ox.ac.uk

**Abstract**—The latest generation of Convolutional Neural Networks (CNN) have achieved impressive results in challenging benchmarks on image recognition and object detection, significantly raising the interest of the community in these methods. Nevertheless, it is still unclear how different CNN methods compare with each other and with previous state-of-the-art shallow representations such as the Bag-of-Visual-Words and the Improved Fisher Vector. This paper conducts a rigorous evaluation of these new techniques, exploring different deep architectures and comparing them on a common ground, identifying and disclosing important implementation details. We identify several useful properties of CNN-based representations, including the fact that the dimensionality of the CNN output layer can be reduced significantly without having an adverse effect on performance. We also identify aspects of deep and shallow methods that can be successfully shared. In particular, we show that the data augmentation techniques commonly applied to CNN-based methods can also be applied to shallow methods, and result in an analogous performance boost. Source code and models to reproduce the experiments in the paper is made publicly available.

### 1 INTRODUCTION

PERHAPS the single most important design choice in current state-of-the-art image classification and object recognition systems is the choice of visual features, or image representation. In fact, most of the quantitative improvements to image

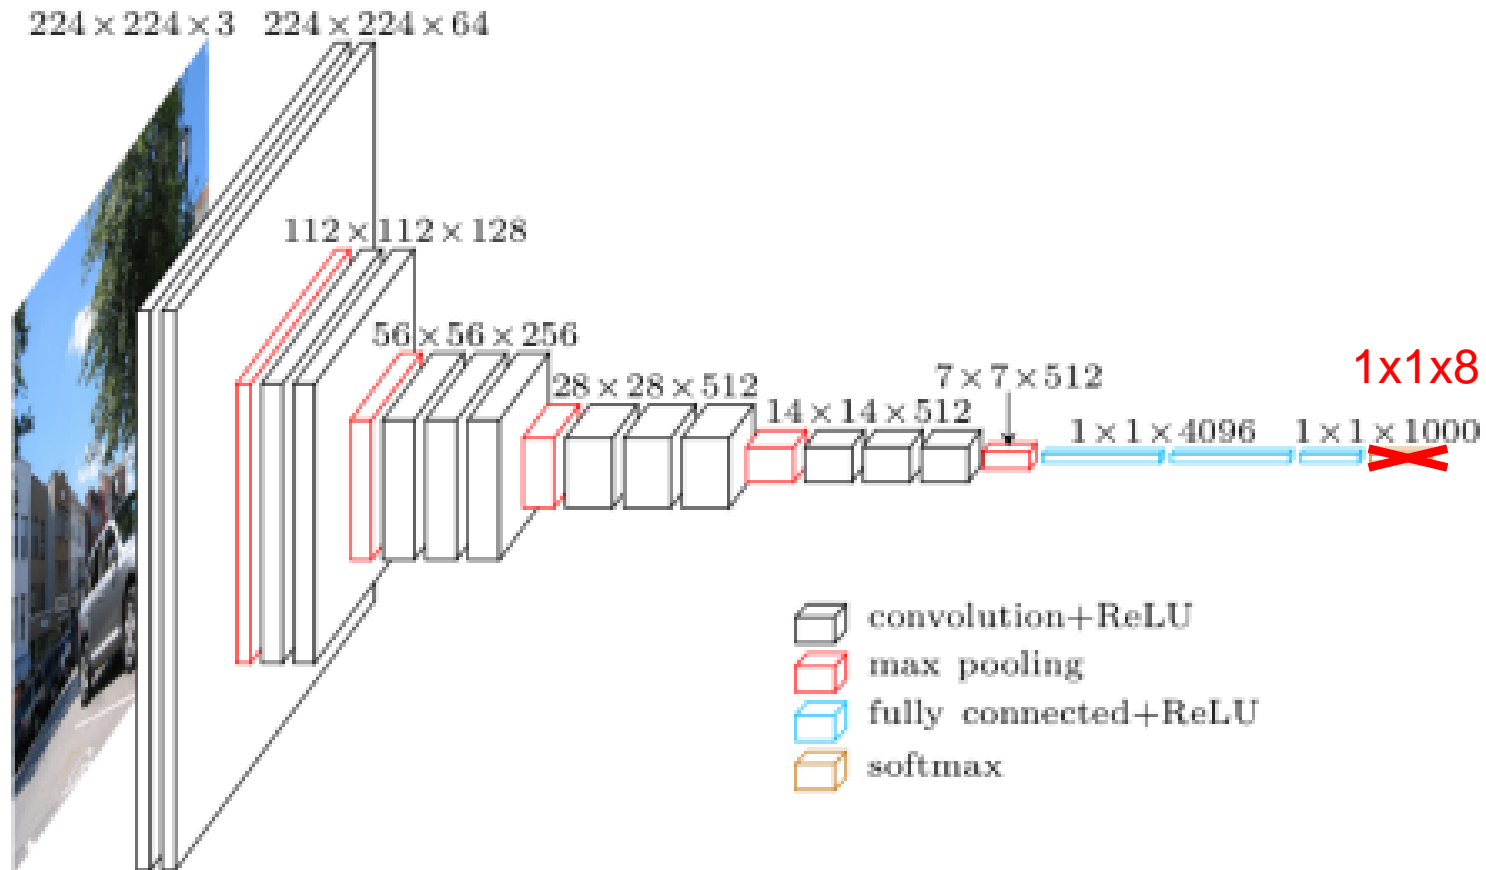
is handcrafted, they contain a very large number of parameters learnt from data. When applied to standard image classification and object detection benchmark datasets such as ImageNet ILSVRC [5] and PASCAL VOC [6] such networks have demonstrated excellent performance [7], [8], [9], [10], [11], significantly better than standard image encod-

## Goals:

- Understand layer manipulation
- Deal with dataset loading
- Hyperparameter optimization

Chatfield, Ken, et al. "Return of the devil in the details: Delving deep into convolutional nets." *arXiv preprint arXiv:1405.3531* (2014).

# Very deep convolutional networks for large-scale image recognition



Credit Davi Frossard

# Understand layer manipulation

input_1: InputLayer	input:	(None, 3, 224, 224)
	output:	(None, 3, 224, 224)



fc1: Dense	input:	(None, 25088)
	output:	(None, 4096)



fc2: Dense	input:	(None, 4096)
	output:	(None, 4096)



predictions: Dense	input:	(None, 4096)
	output:	(None, 1000)

```
img_input = Input(shape=(3,224,224))
```

```
x = Dense(4096, activation='relu', name='fc1')(x)
```

```
x = Dense(4096, activation='relu', name='fc2')(x)
```

```
x = Dense(1000, activation='softmax', name='predictions')(x)
```

```
base_model = Model(img_input, x, name='vgg16')
```

# Understand layer manipulation

input_1: InputLayer	input:	(None, 3, 224, 224)
	output:	(None, 3, 224, 224)

fc1: Dense	input:	(None, 25088)
	output:	(None, 4096)

fc2: Dense	input:	(None, 4096)
	output:	(None, 4096)

predictions: Dense	input:	(None, 4096)
	output:	(None, 1000)

```
img_input = Input(shape=(3,224,224))
```

```
x = Dense(4096, activation='relu', name='fc1')(x)
```

```
x = Dense(4096, activation='relu', name='fc2')(x)
```

~~```
x = Dense(1000, activation='softmax', name='predictions')(x)
```~~

|                    |         |              |
|--------------------|---------|--------------|
| predictions: Dense | input:  | (None, 4096) |
|                    | output: | (None, 8)    |

```
x = base_model.layers[-2].output
```



```
x = Dense(8, activation='softmax', name='predictions')(x)
```

```
model = Model(inputs=base_model.input, outputs=x)
```

# Minimum results

Full training dataset

No hyperparameter optimization



Let's do things more interesting:

- cut the architecture in a lower layer
- use less training data (no more than 400)



## Preparing the model

- Set goal function and process model

```
model.compile(loss='categorical_crossentropy',  
              optimizer='adadelta', metrics=['accuracy'])
```

- Do not train on full network at starting point

```
for layer in base_model.layers:  
    layer.trainable = False
```

# Deal with dataset loading

```
from keras.applications.inception_v3 import preprocess_input

datagen = ImageDataGenerator(featurewise_center=False,
                             samplewise_center=False,
                             featurewise_std_normalization=False,
                             samplewise_std_normalization=False,
                             preprocessing_function=preprocess_input,  IMPORTANT
                             rotation_range=0.,
                             width_shift_range=0.,
                             height_shift_range=0.,
                             shear_range=0.,
                             zoom_range=0.,
                             fill_mode='nearest',
                             horizontal_flip=False,
                             vertical_flip=False,
                             rescale=None)
```

## Deal with dataset loading

```
train_generator = datagen.flow_from_directory(train_data_dir,  
   target_size=(img_width, img_height),  
   batch_size=batch_size,  
   class_mode='categorical')
```

```
test_generator = datagen.flow_from_directory(test_data_dir,  
   target_size=(img_width, img_height),  
   batch_size=batch_size,  
   class_mode='categorical')
```

```
validation_generator = datagen.flow_from_directory(val_data_dir,  
  target_size=(img_width, img_height),  
  batch_size=batch_size,  
  class_mode='categorical')
```

# Deal with dataset loading

```
history=model.fit_generator(train_generator,  
                             samples_per_epoch=400,  
                             nb_epoch=number_of_epoch,  
                             validation_data=validation_generator,  
                             nb_val_samples=800)
```

```
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])
```

```
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])
```

```
result = model.evaluate_generator(test_generator)
```

Afterwards, retrain in full model

# Hyperparameter optimization

Per model

batch\_size = [10, 20, 40, 60, 80, 100]

epochs = [10, 50, 100]

optimizer = ['SGD', 'RMSprop', 'Adagrad', 'Adadelata', 'Adam', 'Adamax', 'Nadam']

learn\_rate = [0.0001, 0.001, 0.01, 0.1, 0.2, 0.3]

momentum = [0.0, 0.2, 0.4, 0.6, 0.8, 0.9]

data augmentation: flip, zoom, rescale, ...

Per layer:

activation = ['softmax', 'softplus', 'softsign', 'relu', 'tanh', 'sigmoid', 'hard\_sigmoid', 'linear']

init\_mode = ['uniform', 'lecun\_uniform', 'normal', 'zero', 'glorot\_normal', 'glorot\_uniform', 'he\_normal', 'he\_uniform'] (Not useful in our case)

Topology:

drop-out layers: p % of inactive weights

batchnormalization

regularizers

# Tasks

Understanding layer manipulation

0. Fine tune an existing architecture

<https://keras.io/applications/>

1. Set a new model from an existing architecture.
2. Apply the model to a small set of data (no more than 400)  
/ghome/mcv/m3/datasets/MIT\_small\_train\_X

Deal with dataset loading

3. Introduce and evaluate the usage of data augmentation

Hyperparameter optimization

4. Introduce and evaluate the usage of any suitable methodology to improve learning curve (dropout layer, batch norm, ...)
5. Apply random search / optuna on per model hyperparameters

|                                   | Size<br>(MB) | Top-1<br>Accuracy | Top-5<br>Accuracy | Parameters | Depth |
|-----------------------------------|--------------|-------------------|-------------------|------------|-------|
| <a href="#">Xception</a>          | 88           | 79.0%             | 94.5%             | 22.9M      | 81    |
| <a href="#">VGG16</a>             | 528          | 71.3%             | 90.1%             | 138.4M     | 16    |
| <a href="#">ResNet50</a>          | 98           | 74.9%             | 92.1%             | 25.6M      | 107   |
| <a href="#">InceptionV3</a>       | 92           | 77.9%             | 93.7%             | 23.9M      | 189   |
| <a href="#">InceptionResNetV2</a> | 215          | 80.3%             | 95.3%             | 55.9M      | 449   |
| <a href="#">MobileNet</a>         | 16           | 70.4%             | 89.5%             | 4.3M       | 55    |
| <a href="#">DenseNet121</a>       | 33           | 75.0%             | 92.3%             | 8.1M       | 242   |
| <a href="#">NASNetMobile</a>      | 23           | 74.4%             | 91.9%             | 5.3M       | 389   |
| <a href="#">EfficientNetB0</a>    | 29           | 77.1%             | 93.3%             | 5.3M       | 132   |
| <a href="#">ConvNeXtTiny</a>      | 109.42       | 81.3%             | -                 | 28.6M      | -     |

## Grades, deliverables and deadline

- Deliver source code and a **short** slide presentation of the work done
  - For each task, all the tests with their associated results
  - 1 slide summarizing the best yielded result and configuration for each task
- Delivered by Monday 30th at 10:30AM



## Control pipeline: Callbacks

- ModelCheckpoint
- EarlyStopping
- ReduceLROnPlateau
- CSVLogger
- LambdaCallback
- ...

Usage:

```
callbacks = [ModelCheckpoint(...), EarlyStopping(...),...]  
model.fit(..., callbacks)
```

# Control pipeline: Callbacks

## Example:

```
plot_loss_callback = LambdaCallback(on_epoch_end=lambda epoch, logs: plot_loss(epoch, logs))
save_callback=ModelCheckpoint(filepath=
                                'weights.{epoch:02d}-{val_loss:.2f}.hdf5',
                                monitor='val_acc', verbose=1, save_best_only=True,
                                mode='max')

history=model.fit(train_generator,
                  samples_per_epoch=1900,
                  nb_epoch=number_of_epoch,
                  validation_data=validation_generator,
                  nb_val_samples=800,
                  callbacks=[plot_loss_callback, save_callback])
```