


Deep Learning for Object Detection

Lluís Gomez i Bigorda

Object detection pipeline (recap of previous class)


Given the unbalanced nature of detection. What do we need?

Two Stage Framework




Object Detection Models (recap of previous class)

OverFeat




R-CNN




Object Detection Models (recap of previous class)

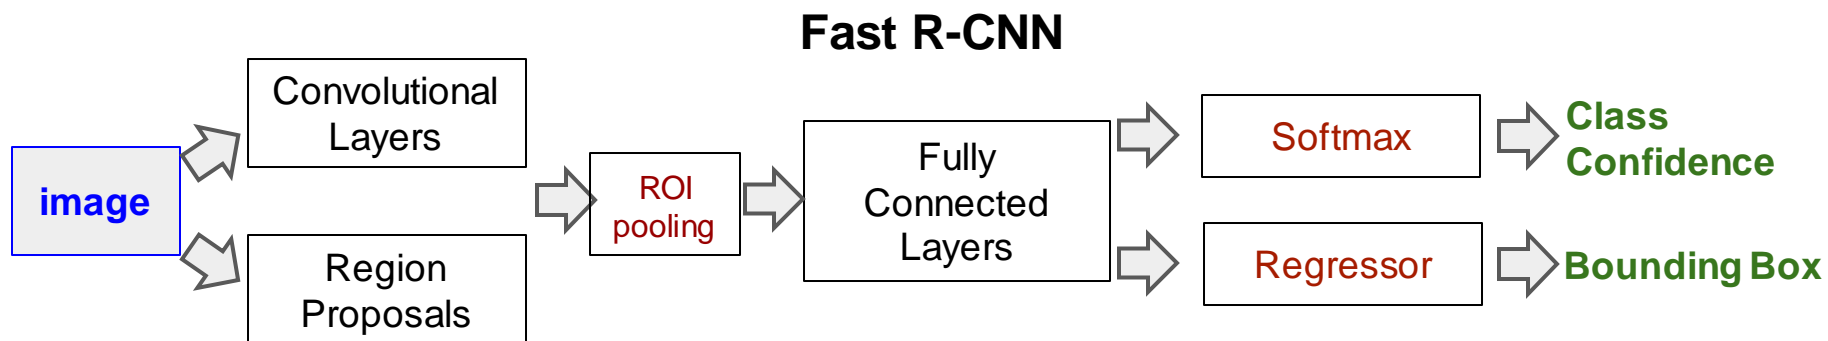

R-CNN



SPP




Object Detection Models (recap of previous class)

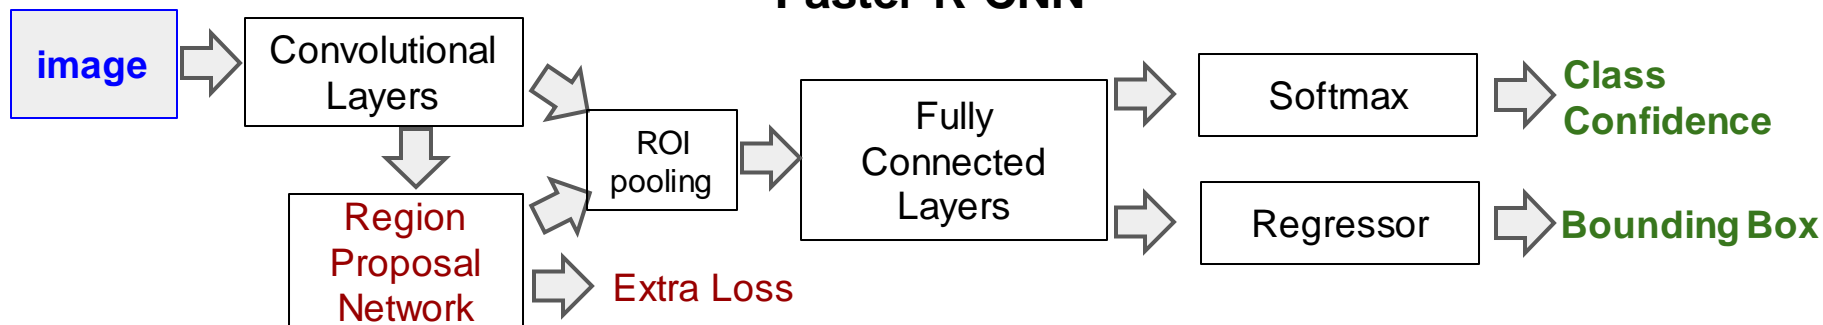


Object Detection Models (recap of previous class)

Fast R-CNN



Faster R-CNN



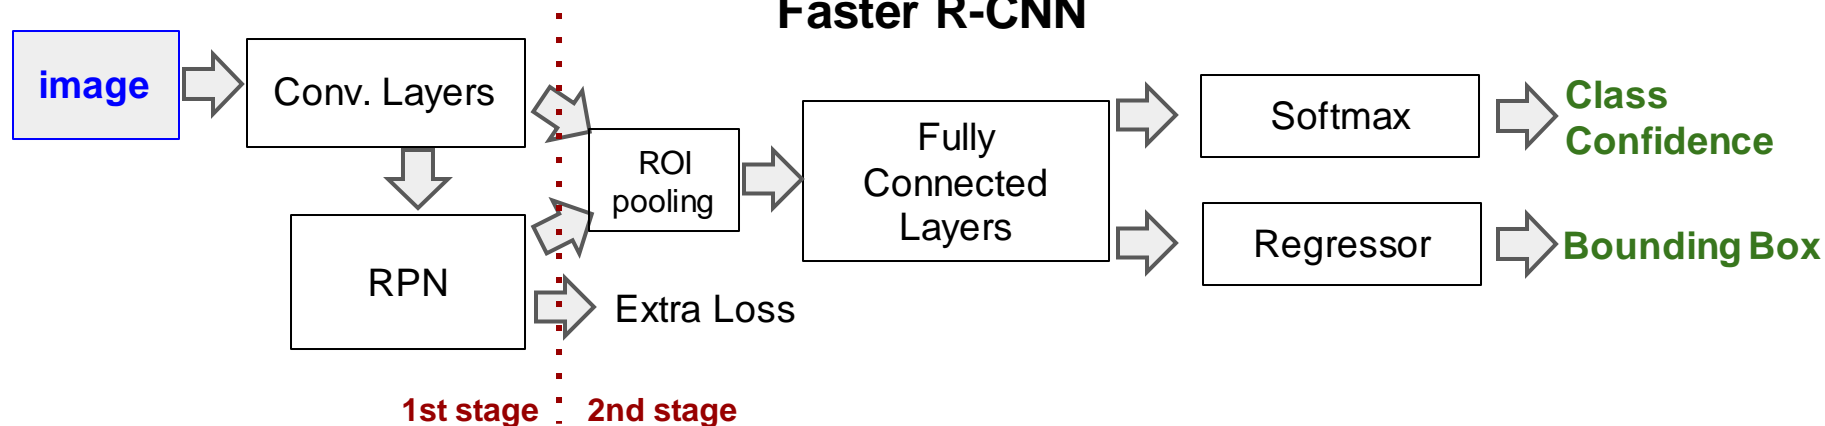
Deep learning for object detection: Outline

- Introduction
- Basic blocks and concepts
- Models (i)
- Models (ii)
 - Single Stage Object Detectors
 - Feature Pyramid Networks
 - Focal Loss
 - Mask R-CNN
 - DETR
 - Other ideas

You Only Look Once - YOLO (2015, CVPR2016)


Two Stage Framework

Faster R-CNN




Single Stage Framework


YOLO




YOLO: Key idea



YOLO: Architecture



YOLO: Training




1. Pre-train network on Imagenet classification task
2. Train the model with joint loss (quite engineered loss function)

YOLO: Training tricks

1. Use 448×448 input for detection, instead of 224×224 ,
2. Use Leaky ReLU for all layers,
3. Dropout after the first fully connected layer,
4. Normalize bounding boxes parameters in $[0, 1]$,
5. Use a quadratic loss not only for the bounding box coordinates, but also for the confidence and the class scores,
 1. Reduce the weight of large bounding boxes by using the square roots of the size in the loss,
 2. Reduce the importance of empty cells by weighting less the confidence-related loss on them,
1. Use momentum 0.9, decay $5e - 4$,
2. Data augmentation with scaling, translation, and HSV transformation.

YOLO: Inference





Single pass through the network.

Inference is very fast.

https://docs.google.com/presentation/d/1kAa7NOamBt4calBU9iHgT8a86RRHz9Yz2oh4-GTdX6M/edit#slide=id.g151008b386_0_0


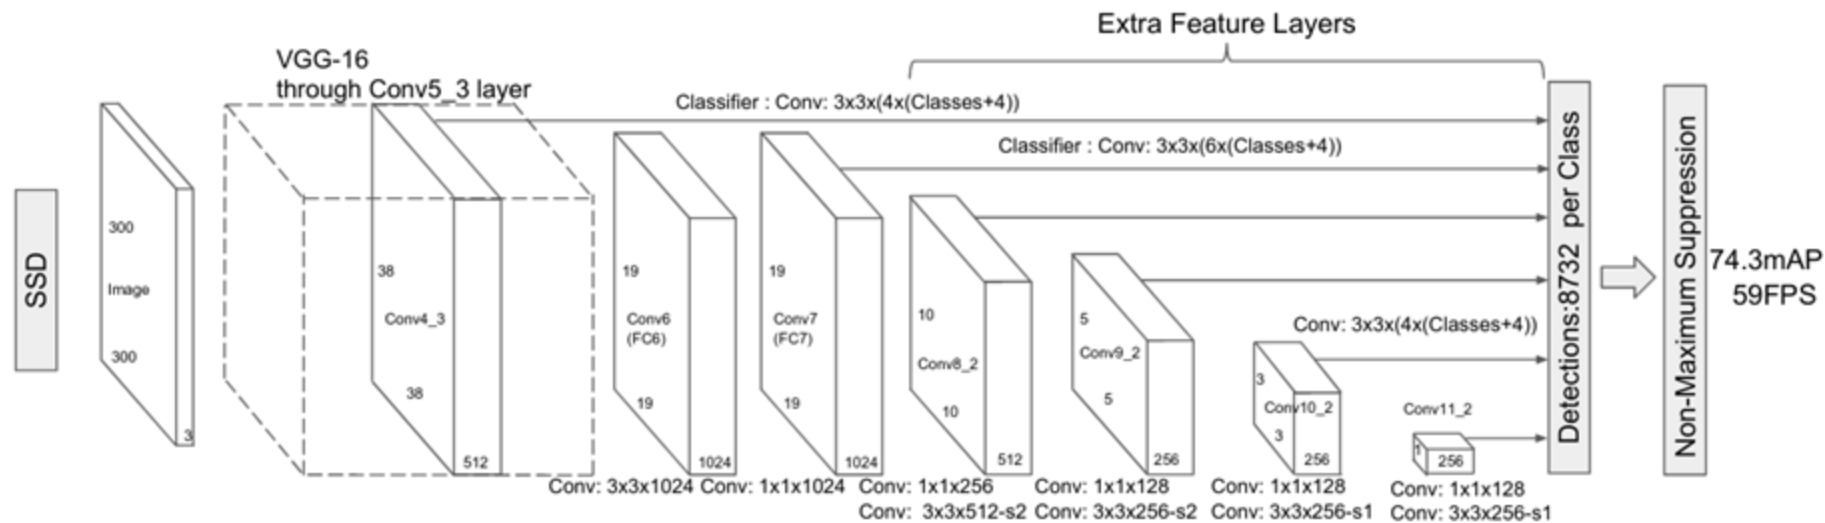
YOLO: Results

	Pascal 2007 mAP	Speed	
DPM v5	33.7	.07 FPS	14 s/img
R-CNN	66.0	.05 FPS	20 s/img
Fast R-CNN	70.0	.5 FPS	2 s/img
Faster R-CNN	73.2	7 FPS	140 ms/img
YOLO	63.4	45 FPS	22 ms/img




https://docs.google.com/presentation/d/14qBAiyhMOFI_wZW4dA1CkixgXwf0zKGbpw_0oHK8yEM/edit?usp=sharing

Single Shot Detector / SSD (ECCV 2016)




LIU, Wei, et al. [SSD: Single shot multibox detector](#). ECCV 2016.


Single Shot Detector / SSD (ECCV 2016)



(a) Image with GT boxes



(b) 8×8 feature map




(c) 4×4 feature map

YOLOv2 (2016)

	YOLO									YOLOv2
batch norm?		✓	✓	✓	✓	✓	✓	✓	✓	✓
hi-res classifier?			✓	✓	✓	✓	✓	✓	✓	✓
convolutional?				✓	✓	✓	✓	✓	✓	✓
anchor boxes?				✓	✓					
new network?					✓	✓	✓	✓	✓	✓
dimension priors?						✓	✓	✓	✓	✓
location prediction?						✓	✓	✓	✓	✓
passthrough?							✓	✓	✓	✓
multi-scale?								✓	✓	✓
hi-res detector?									✓	✓
VOC2007 mAP	63.4	65.8	69.5	69.2	69.6	74.4	75.4	76.8		78.6

There are a lot of tricks to get a good architecture for object detection...

Comparison



The YOLO family from v1 to v7

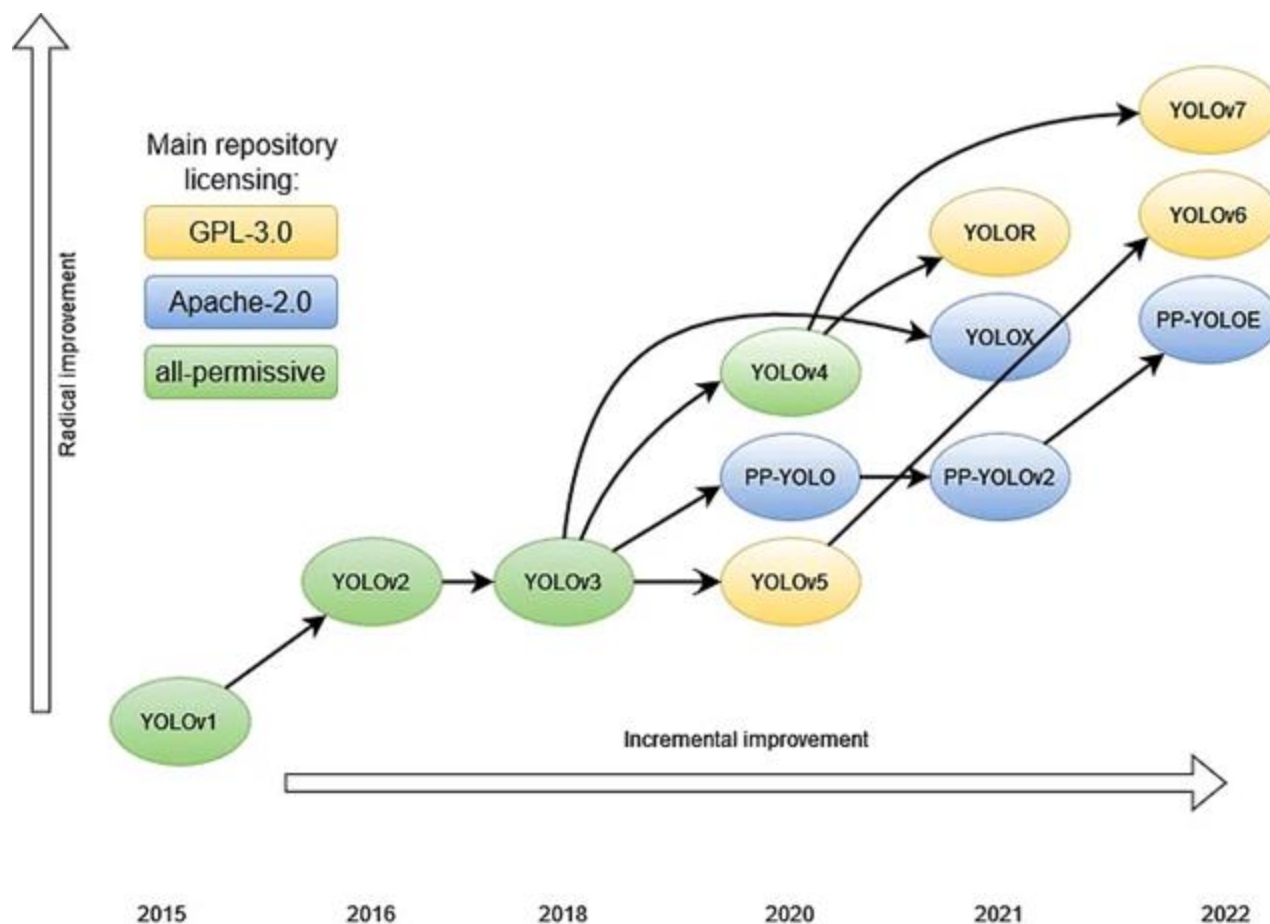


Figure source: [The evolution of the YOLO neural networks family from v1 to v7.](#)

The YOLO family from v1 to v7

YOLO version	backbone	neck	head(s)	augmentations
YOLOv1	GoogLeNet, VGG-16	2x fully connected layers	combined classes + bboxes	random scaling & translations up to 20%; random adjust exposure & saturation up to x1.5 in HSV
YOLOv2	Darknet-19	fully convolutional layers	combined hierarchical classes + bboxes, anchor-based	random crops, rotations, and hue, saturation, and exposure shifts
YOLOv3	Darknet-53	FPN	combined multilabel + bboxes, anchor-based	no specific info, seems like the same as in YOLOv2
YOLOv5	CSPDarknet53	SPPF, CSP-PAN	combined multilabel + bboxes, anchor-based	Mosaic, copy-paste, random affine, MixUp, random adjust HSV, random horizontal flip
PP-YOLO	ResNet50-vd + deformable convolutions	FPN, SPP	combined multilabel + bboxes, anchor-based	MixUp
YOLOv4	CSPDarknet53	PANet, SPP	combined multilabel + bboxes, anchor-based	CutMix, Mosaic, MixUp, CutOut, Self-Adversarial Training, bilateral blurring
PP-YOLOv2	ResNet50-vd + deformable convolutions	PANet	combined multilabel + bboxes, anchor-based	MixUp; random color distortion, expand, crop, flip
YOLOX	Darknet-53	FPN	decoupled multilabel + bboxes, anchor-free	Mosaic, MixUp, random horizontal flip, colorjitter
YOLOR	sequence of convolutional layers with downscaling	FPN, CSP, SPP	multi-head (object detection, multi-label classification, feature embedding)	CutMix, Mosaic, MixUp, CutOut, Self-Adversarial Training, bilateral blurring
PP-YOLOE	CSPRepResNet	PANet	Efficient Task-aligned Head (decoupled), anchor-free	random crop, horizontal flip, color distortion, multi-scale
YOLOv6	EfficientRep	Rep-PAN	Efficient decoupled head, anchor-free	Mosaic, MixUp
YOLOv7	Extended-ELAN	-	multiple (lead heads & aux heads), anchor-based	random perspective, HSV jitter, flips, Mosaic

Table source: [The evolution of the YOLO neural networks family from v1 to v7.](#)

Non-Maximum Suppression (NMS)

(remember) Common component to all Object Detection architectures!




Image
Predictions



- 1) Order bb by confidence
- 2) Pick the most confident bb
- 3) Remove all bb with $IoU > th$





Image
Predictions

Non-Maximum Suppression (NMS)

Before non-max suppression



Non-Max
Suppression




After non-max suppression




<https://towardsdatascience.com/non-maximum-suppression-nms-93ce178e177c>

Non-Maximum Suppression (NMS)



<https://towardsdatascience.com/non-maximum-suppression-nms-93ce178e177c>

Non-Maximum Suppression (NMS)



Bodla et al. [Soft-NMS -- Improving Object Detection With One Line of Code](#) ICCV 2017

Non-Maximum Suppression (NMS)

Input : $\mathcal{B} = \{b_1, \dots, b_N\}$, $\mathcal{S} = \{s_1, \dots, s_N\}$, N_t
 \mathcal{B} is the list of initial detection boxes
 \mathcal{S} contains corresponding detection scores
 N_t is the NMS threshold

```
begin
   $\mathcal{D} \leftarrow \{\}$ 
  while  $\mathcal{B} \neq \text{empty}$  do
     $m \leftarrow \operatorname{argmax} \mathcal{S}$ 
     $\mathcal{M} \leftarrow b_m$ 
     $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{M}$ ;  $\mathcal{B} \leftarrow \mathcal{B} - \mathcal{M}$ 
    for  $b_i$  in  $\mathcal{B}$  do
      if  $\operatorname{iou}(\mathcal{M}, b_i) \geq N_t$  then
         $\mathcal{B} \leftarrow \mathcal{B} - b_i$ ;  $\mathcal{S} \leftarrow \mathcal{S} - s_i$ 
      end
    end
  end
  return  $\mathcal{D}, \mathcal{S}$ 
end
```

NMS

Soft-NMS

Non-Maximum Suppression (NMS)




NMS : <https://github.com/rbgirshick/fast-rcnn/blob/master/lib/utils/nms.py>




Soft-NMS : https://github.com/DocF/Soft-NMS/blob/master/soft_nms.py


Feature Pyramid Networks (CVPR 2017)




(a) Featurized image pyramid



(b) Single feature map




(c) Pyramidal feature hierarchy



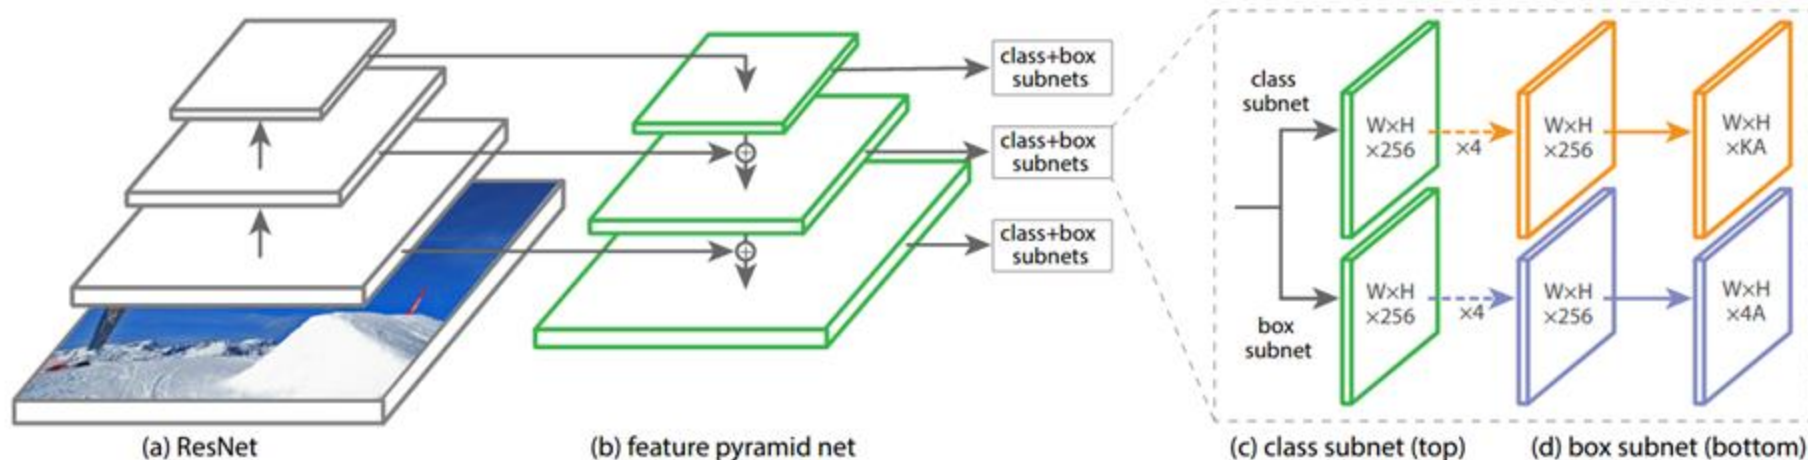
(d) Feature Pyramid Network

Feature Pyramid Networks



Faster R-CNN	proposals	feature	head	lateral?	top-down?	AP@0.5	AP	AP _s	AP _m	AP _l
(*) baseline from He <i>et al.</i> [16] [†]	RPN, C_4	C_4	conv5			47.3	26.3	-	-	-
(a) baseline on conv4	RPN, C_4	C_4	conv5			53.1	31.6	13.2	35.6	47.1
(b) baseline on conv5	RPN, C_5	C_5	2fc			51.7	28.0	9.6	31.9	43.1
(c) FPN	RPN, $\{P_k\}$	$\{P_k\}$	2fc	✓	✓	56.9	33.9	17.8	37.7	45.8

RetinaNet (ICCV 2017)



4 Conv Layers with 256 3x3 filters

$A=9$ **anchor boxes**


$K=80$ object class labels (COCO)

Focal Loss

$$\text{CE}(p_t) = -\log(p_t)$$

$$\text{FL}(p_t) = -(1 - p_t)^\gamma \log(p_t)$$

$$p_t = \begin{cases} p & \text{if } y = 1 \\ 1 - p & \text{otherwise,} \end{cases}$$



The unbalanced nature of detection. Hard Negative Mining vs. Focal Loss


Why putting more focus on hard, misclassified examples?



Positive	1
Negative	70
Hard negative	6

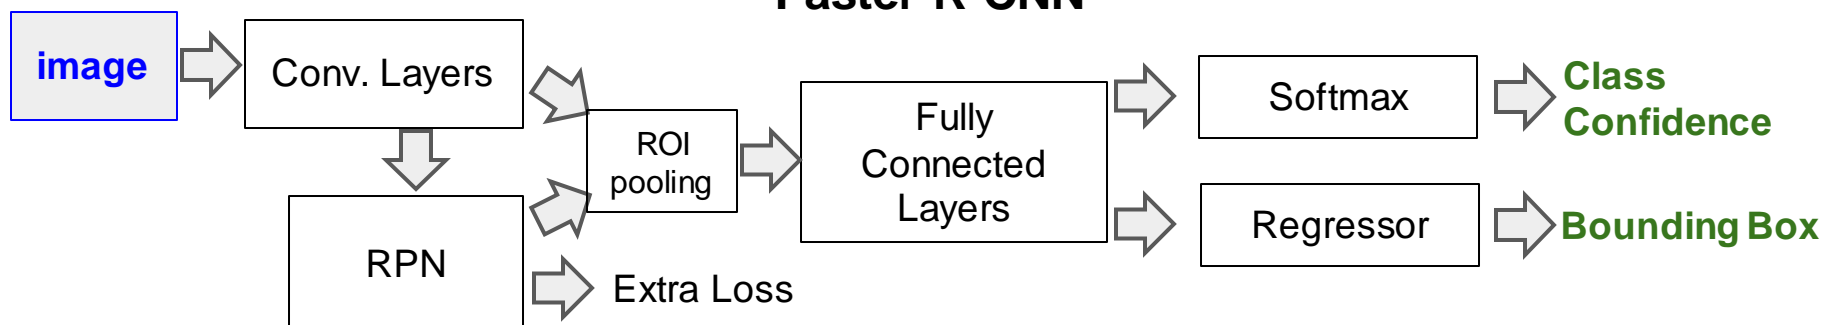
<https://medium.com/@ageitgey/machine-learning-is-fun-part-3-deep-learning-and-convolutional-neural-networks-f40359318721#.dnbyjd6zg>

Focal Loss / RetinaNet




Mask R-CNN (ICCV 2017)


Faster R-CNN



Mask R-CNN




Mask R-CNN for instance segmentation




He, K., Gkioxari, G., Dollár, P., & Girshick, [R. Mask R-CNN](#). ICCV 2017.


ROI Align




Input activation



Region projection and pooling sections



Sampling locations





Bilinear interpolated values

2×2 values per cell.



Max pooling output

ROI Align / Bilinear interpolation



Bilinear interpolation for RoIAlign.

Mask R-CNN Bounding Box Detection Results

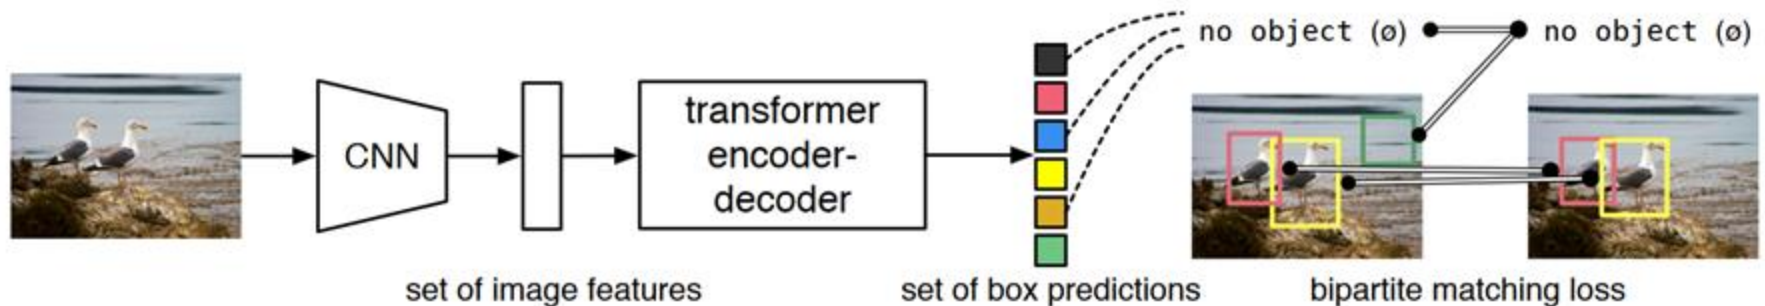
	backbone	AP^{bb}	AP_{50}^{bb}	AP_{75}^{bb}	AP_S^{bb}	AP_M^{bb}	AP_L^{bb}
Faster R-CNN+++ [19]	ResNet-101-C4	34.9	55.7	37.4	15.6	38.7	50.9
Faster R-CNN w FPN [27]	ResNet-101-FPN	36.2	59.1	39.0	18.2	39.0	48.2
Faster R-CNN by G-RMI [21]	Inception-ResNet-v2 [41]	34.7	55.5	36.7	13.5	38.1	52.0
Faster R-CNN w TDM [39]	Inception-ResNet-v2-TDM	36.8	57.7	39.2	16.2	39.8	52.1
Faster R-CNN, RoIAlign	ResNet-101-FPN	37.3	59.6	40.3	19.8	40.2	48.8
Mask R-CNN	ResNet-101-FPN	38.2	60.3	41.7	20.1	41.1	50.2
Mask R-CNN	ResNeXt-101-FPN	39.8	62.3	43.4	22.1	43.2	51.2

He, K., Gkioxari, G., Dollár, P., & Girshick, [R. Mask R-CNN](#). ICCV 2017.

DETR (ECCV 2020)

End-to-End Object Detection with Transformers


- DETR directly predicts (in parallel) the final set of detections by combining a CNN with a transformer architecture. No need of NMS!
- During training, bipartite matching uniquely assigns predictions with ground truth boxes. Prediction with no match should yield a “no object” (\emptyset) class prediction.



DETR (ECCV 2020)

End-to-End Object Detection with Transformers

- Encoder input: CNN features + positional encoding.
- Decoder input:
 - a fixed number N of learned positional embeddings ($N=100$), **object queries**.
 - also attends to the encoder output.
- Output embeddings of the decoder go to a shared feed forward network (FFN) that predicts a detection (class and bbox) or a “no object” class.



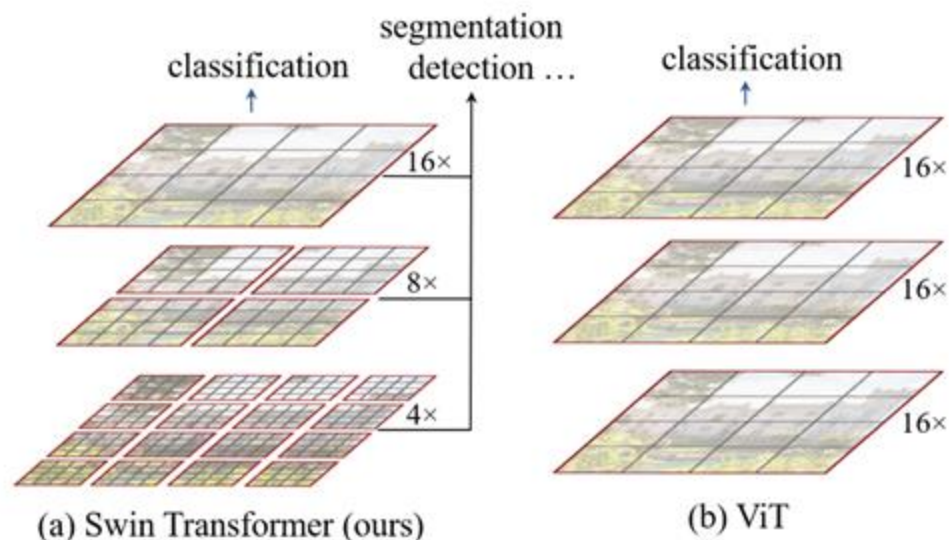
DETR (ECCV 2020)

Model	GFLOPS/FPS	#params	AP	AP ₅₀	AP ₇₅	AP _S	AP _M	AP _L
Faster RCNN-DC5	320/16	166M	39.0	60.5	42.3	21.4	43.5	52.5
Faster RCNN-FPN	180/26	42M	40.2	61.0	43.8	24.2	43.5	52.0
Faster RCNN-R101-FPN	246/20	60M	42.0	62.5	45.9	25.2	45.6	54.6
Faster RCNN-DC5+	320/16	166M	41.1	61.4	44.3	22.9	45.9	55.0
Faster RCNN-FPN+	180/26	42M	42.0	62.1	45.5	26.6	45.4	53.4
Faster RCNN-R101-FPN+	246/20	60M	44.0	63.9	47.8	27.2	48.1	56.0
DETR	86/28	41M	42.0	62.4	44.2	20.5	45.8	61.1
DETR-DC5	187/12	41M	43.3	63.1	45.9	22.5	47.3	61.1
DETR-R101	152/20	60M	43.5	63.8	46.4	21.9	48.0	61.8
DETR-DC5-R101	253/10	60M	44.9	64.7	47.7	23.7	49.5	62.3

Swin Transformer (2021)

The hierarchical Transformers (e.g., Swin Transformers) reintroduce several ConvNet priors.

Makes Transformers practically viable as a generic vision backbone and demonstrate remarkable performance on a wide variety of vision tasks.



Swin Transformer (2021)

(b) Various backbones w. Cascade Mask R-CNN									
	AP^{box}	AP_{50}^{box}	AP_{75}^{box}	AP^{mask}	AP_{50}^{mask}	AP_{75}^{mask}	param	FLOPs	FPS
DeiT-S [†]	48.0	67.2	51.7	41.4	64.2	44.3	80M	889G	10.4
R50	46.3	64.3	50.5	40.1	61.7	43.4	82M	739G	18.0
Swin-T	50.5	69.3	54.9	43.7	66.6	47.1	86M	745G	15.3
X101-32	48.1	66.5	52.4	41.6	63.9	45.2	101M	819G	12.8
Swin-S	51.8	70.4	56.3	44.7	67.9	48.5	107M	838G	12.0
X101-64	48.3	66.4	52.3	41.7	64.0	45.1	140M	972G	10.4
Swin-B	51.9	70.9	56.5	45.0	68.4	48.7	145M	982G	11.6

ConvNext (2022)


“Modernize” a standard ResNet toward the design of a vision Transformer: “Patchify” input, ResNeXt, Larger Kernels, ReLU->GeLU, fewer activations, BN->LN, modern optimizer, better augmentations, etc...

Constructed entirely from standard ConvNet modules.

ConvNeXts compete favorably with Transformers in terms of accuracy and scalability, outperforming Swin Transformers on COCO detection, while maintaining the simplicity and efficiency of standard ConvNets.

backbone	FLOPs	FPS	AP ^{box}	AP ₅₀ ^{box}	AP ₇₅ ^{box}	A
Mask-RCNN 3× schedule						
○ Swin-T	267G	23.1	46.0	68.1	50.3	
● ConvNeXt-T	262G	25.6	46.2	67.9	50.8	
Cascade Mask-RCNN 3× schedule						
● ResNet-50	739G	16.2	46.3	64.3	50.5	
● X101-32	819G	13.8	48.1	66.5	52.4	
● X101-64	972G	12.6	48.3	66.4	52.3	
○ Swin-T	745G	12.2	50.4	69.2	54.7	
● ConvNeXt-T	741G	13.5	50.4	69.1	54.8	
○ Swin-S	838G	11.4	51.9	70.7	56.3	
● ConvNeXt-S	827G	12.0	51.9	70.8	56.5	
○ Swin-B	982G	10.7	51.9	70.5	56.4	
● ConvNeXt-B	964G	11.4	52.7	71.3	57.2	
○ Swin-B [‡]	982G	10.7	53.0	71.8	57.5	
● ConvNeXt-B [‡]	964G	11.5	54.0	73.1	58.8	
○ Swin-L [‡]	1382G	9.2	53.9	72.4	58.8	
● ConvNeXt-L [‡]	1354G	10.0	54.8	73.8	59.8	
● ConvNeXt-XL [‡]	1898G	8.6	55.2	74.2	59.9	


Other ideas in Object Detection: Rotated objects



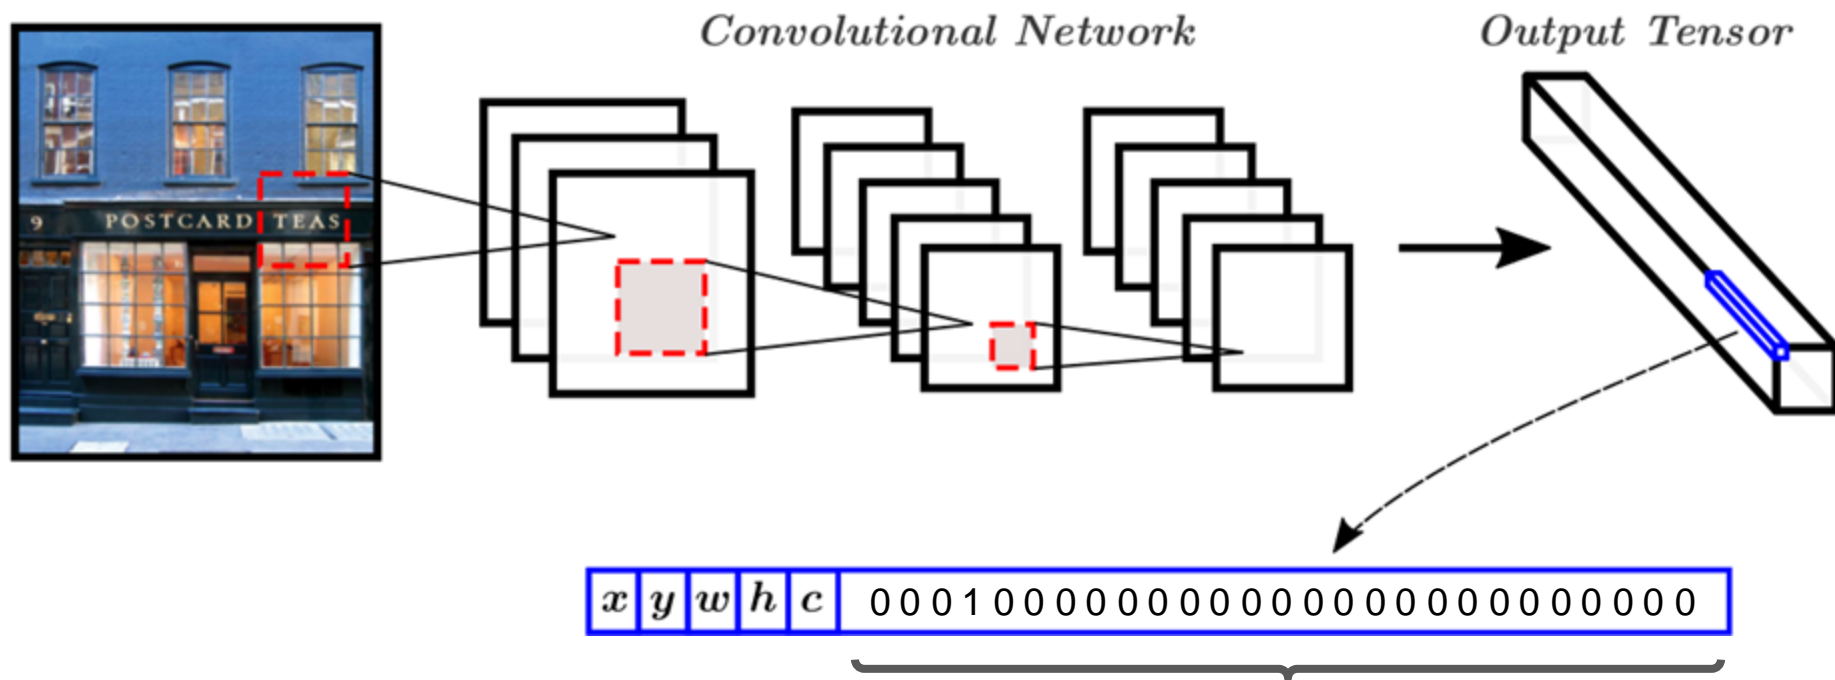
(a)



(b)



Other ideas in Object Detection: STR




One-hot classification

How many classes?

Label Embedding (PHOC)


Text strings are embedded into a d-dimensional binary space: Pyramidal Histogram Of Characters (PHOC)

PHOC encodes if a particular character appears in a particular region of the string



Label Embeddings

Attribute-based recognition



Comparison of:

- **Direct Attribute Prediction (DAP):** compute attribute probabilities + combine scores

Lampert, Nickisch, Harmeling, "Learning To Detect Unseen Object Classes by Between-Class Attribute Transfer", CVPR'09

- **Attribute Label Embedding (ALE):** embed classes + bilinear compatibility

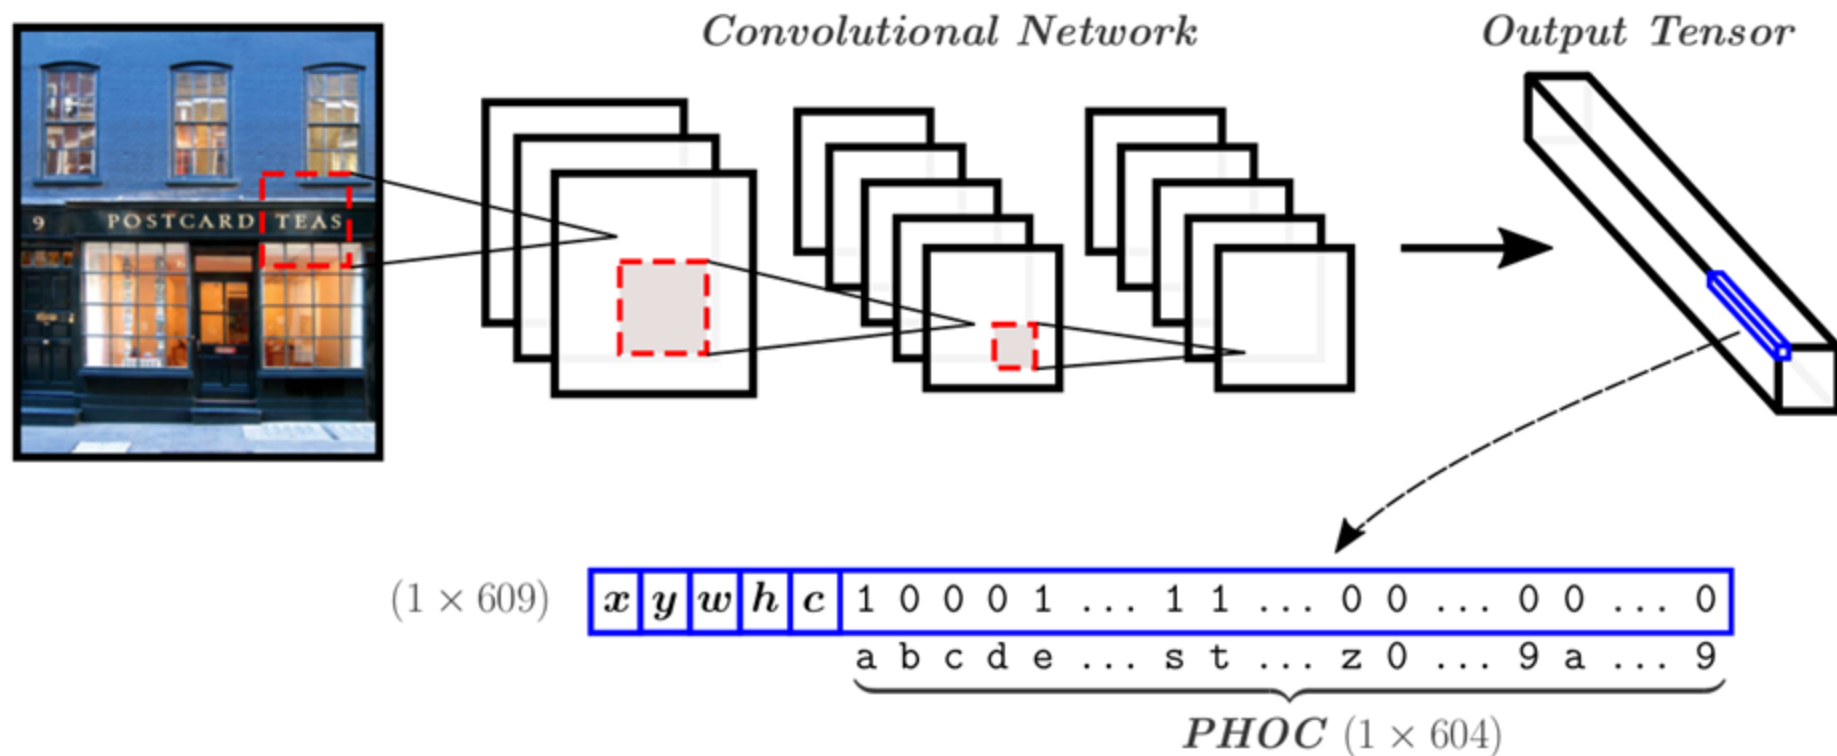
Akata, Perronnin, Harchaoui, Schmid, "Label-embedding for attribute-based classification", CVPR'13

→ **ALE outperforms DAP by large margin on zero-shot bird recognition**

See also: Alabdulmohsin, Cissé, Zhang, "Is attribute-based zero-shot learning an ill-posed strategy?", EACL'16.

Slide credit: Florent Perronnin. "Output Embedding for Large Scale Computer Vision", ECCV 2018.

Single Shot Text Detection and Recognition



L.Gomez, A. Mafla, M. Rusiñol, D. Karatzas. "Single Shot Scene Text Retrieval", ECCV 2018.

References

- Ross B. Girshick, Jeff Donahue, Trevor Darrell and Jitendra Malik; Rich feature hierarchies for accurate object detection and semantic segmentation.
- Pierre Sermanet, David Eigen, Xiang Zhang, Michael Mathieu, Rob Fergus, Yann LeCun; OverFeat: Integrated Recognition, Localization and Detection using Convolutional Networks.
- Kaiming He and Xiangyu Zhang and Shaoqing Ren and Jian Sun; Spatial Pyramid Pooling in Deep Convolutional Networks for Visual Recognition.
- Ross B. Girshick; Fast R-CNN.
- Shaoqing Ren, Kaiming He, Ross B. Girshick and Jian Sun; Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks.
- Joseph Redmon, Santosh Kumar Divvala, Ross B. Girshick, Ali Farhadi; You Only Look Once: Unified, Real-Time Object Detection.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun; Deep Residual Learning for Image Recognition.
- Joseph Redmon and Ali Farhadi; YOLO9000: Better, Faster, Stronger.
- M. Oquab and L. Bottou and I. Laptev and J. Sivic; Is object localization for free? - Weakly-supervised learning with convolutional neural networks.

References

- Liu, W., Anguelov, D., Erhan, D., Szegedy, C., Reed, S., Fu, C. Y., & Berg, A. C. (2016, October). SSD: Single shot multibox detector. ECCV 2016.
- Bodla, N., Singh, B., Chellappa, R., & Davis, L. S. Soft-NMS--improving object detection with one line of code. ICCV 2017.
- Lin, T. Y., Dollár, P., Girshick, R., He, K., Hariharan, B., & Belongie, S. (2017). Feature pyramid networks for object detection. CVPR 2017.
- Lin, T. Y., Goyal, P., Girshick, R., He, K., & Dollár, P. Focal loss for dense object detection. ICCV 2017.
- He, K., Gkioxari, G., Dollár, P., & Girshick, R. Mask R-CNN. ICCV 2017.
- Carion et al."End-to-end object detection with transformers. ECCV, 2020.
- Zhou, X., Yao, C., Wen, H., Wang, Y., Zhou, S., He, W., & Liang, J. EAST: an efficient and accurate scene text detector. CVPR 2017.
- L.Gomez, A. Mafla, M. Rusiñol, D. Karatzas. Single Shot Scene Text Retrieval, ECCV 2018.