



Master in Computer Vision *Barcelona*

Module: Video Analysis

Lecture 5: Bayesian tracking (I)

Lecturer: Ramon Morros

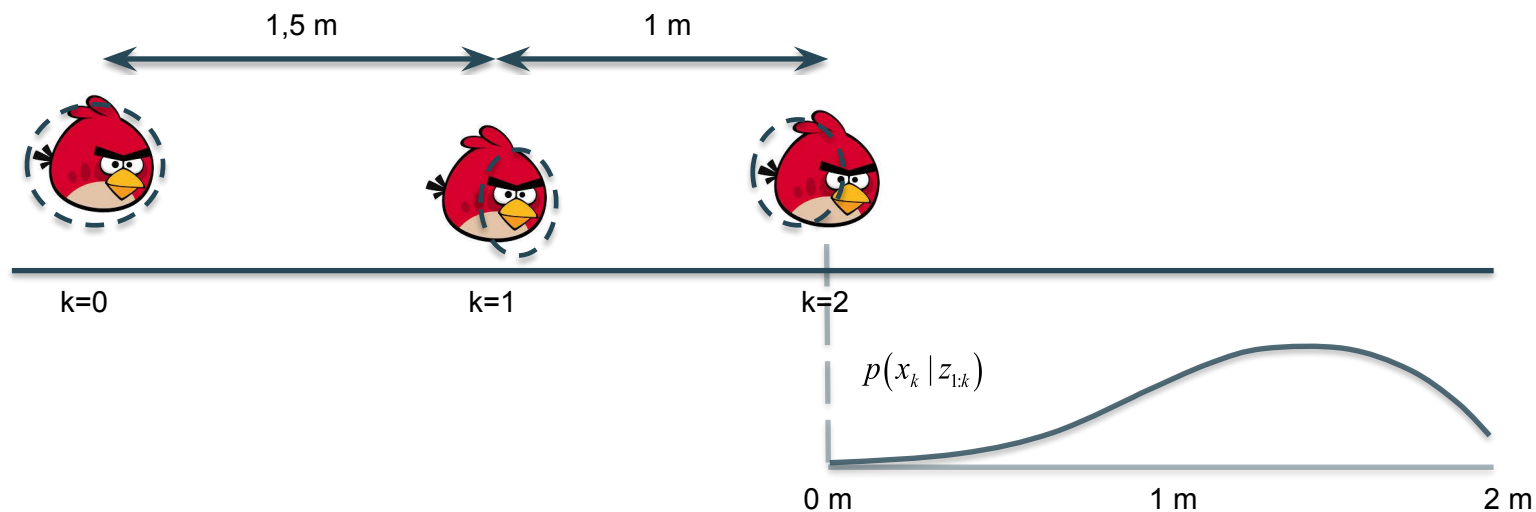
Overview

- Introduction to the tracking problem
- Linear Dynamic Models
- Kalman Filter
- Particle Filters

Tracking

Bayesian estimation

- **Goal:** finding an **object position** over time in a video sequence.



- **How to estimate this function iteratively?**
 - Kalman Filter
 - Particle Filter

Tracking

Observations and states


- **State (x):** true parameters that characterize the object being tracked
 - Position, velocity, etc.
 - Usually “hidden” or unknown
- **Measurement/observation (z):** what we can measure on the image at a given time
 - Blob centroids, bounding boxes, etc.
 - Observations result from underlying states
 - Subject to noise



Tracking


Tracking as inference

- At each time step, **object state changes** (from x_{k-1} to x_k) and we get a new observation z_k
(k indicates discrete time)
- Goal:** recover most likely state x_k given:
 - All observations seen so far (z_1, \dots, z_k)
 - Knowledge about dynamics of state transitions.



Tracking

Steps of tracking



Tracking


Steps of tracking

- **Prediction:** What is the current state of the object given past measurements?

$$p(x_k | z_1, \dots, z_{k-1}) = p(x_k | z_{1:k-1})$$

- **Update:** Compute an updated estimate of the state from prediction and measurements.

$$p(x_k | z_1, \dots, z_{k-1}, z_k) = p(x_k | z_{1:k})$$




Tracking

Questions

- How to represent the dynamics model that govern the changes in the states?
- How to represent the relationship between state and measurements, plus our uncertainty in the measurements?
- How to compute each cycle of updates?
- How to combine prediction and correction?
 - If the dynamics model is too strong, will end up ignoring the data
 - If the observation model is too strong tracking is the observation model is too strong, tracking is reduced to repeated detection

Slide credit: Kristen Grauman


Bayesian filtering



Introduction

Kalman Filter

- Definition:** Algorithm that processes measurements to deduce the **optimal estimate** of the state of a **linear system** using a set of measurements and a **statistical model** of the system.



$$\begin{aligned}\phi(x_{k-1}, v_{k-1}) &= Dx_{k-1} + \Sigma_d \\ \psi(x_{t-1}, v_{t-1}) &= Mx_{t-1} + \Sigma_m\end{aligned}$$



$$\begin{aligned}\mathbf{x}_t &\sim N(\mathbf{D}\mathbf{x}_{t-1}; \Sigma_d) \\ \mathbf{z}_t &\sim N(\mathbf{M}\mathbf{x}_t; \Sigma_m)\end{aligned}$$



$$P(X_t | z_0, \dots, z_{t-1})$$

μ_t^-, σ_t^-
Mean and std. dev. of predicted state

Prediction

Receive measurement

Update

Time advances: $t++$

$$P(X_t | z_0, \dots, z_t)$$

μ_t^+, σ_t^+
Mean and std. dev. of predicted state

Tracking

Assumptions

- **Markov process:** Current state X_k depends **only** on previous state X_{k-1}

$$p(x_k | x_0, \dots, x_{k-1}) = p(x_k | x_{k-1})$$

Dynamics model

- **Independence:** Measurement depends only on current state

$$p(z_t | x_0, \dots, x_k) = p(z_k | x_k)$$

Observation model

Tracking

Filtering framework

- Discrete-time state space filtering
- We want to **recursively** estimate the current state at every time that a measurement is received
- **Prediction:**
 - Propagate state pdf forward in time, taking noise into account (translate, deform, and spread the pdf)
- **Update:**
 - Use Bayes theorem to modify prediction pdf based on current measurement

pdf: the probability density function of a continuous random variable is a function that describes the relative likelihood for this random variable to take on a given value.

Tracking

Bayesian estimation

- Consider

$\{x_k\}_{k \geq 0}$: **Unobserved** process with transition density $x_k | x_{k-1} \sim f(\cdot | x_{k-1})$

$\{z_k\}_{k \geq 1}$: **Observations**, conditionally independent given $\{x_k\}_{k \geq 0}$, of marginal density $z_k | x_k \sim g(\cdot | x_k)$

$$\left[\begin{array}{ll} x_k = \phi(x_{k-1}, v_{k-1}) & \phi, \psi : \text{Deterministic mappings} \\ z_k = \psi(x_k, w_k) & \{v_k\}_{k \geq 2}, \{w_k\}_{k \geq 1} : \text{Independent noises} \end{array} \right.$$

Tracking

Tracking as induction

- Goal: Our goal is to obtain $p(x_k | z_{1:k})$
- Base case:
 - Assume we have an initial **prior** that predicts state in absence of any evidence: $p(x_0)$
 - At the first frame, **update** given the value of z_0

$$\underbrace{p(x_0 | z_0)}_{\text{Posterior prob. of state given measurement}} = \frac{p(z_0 | x_0)p(x_0)}{p(z_0)} \propto \underbrace{p(z_0 | x_0)}_{\text{Likelihood of measurement}} \underbrace{p(x_0)}_{\text{Prior of the state}}$$

Tracking

Tracking as induction

- Prediction:

(Chapman-Kolmogorov equation)

$$p(x_k | z_{1:k-1}) = \int \underbrace{p(x_k | x_{k-1})}_{\text{Dynamics model}} \underbrace{p(x_{k-1} | z_{1:k-1})}_{\text{Corrected estimate from previous step}} dx_{k-1}$$

- Update:

$$p(x_k | z_{1:k}) = \frac{\underbrace{p(z_k | x_k)}_{\text{Observation model}} \underbrace{p(x_k | z_{1:k-1})}_{\text{Predicted estimate}}}{\int p(z_k | x_k) p(x_k | z_{1:k-1}) dx_k}$$

Slide credit: Svetlana Lazebnik

LINEAR DYNAMIC MODELS

Linear Dynamics Models (LDM)

Linear dynamics models

- Dynamics model

- State undergoes linear transformation **D** plus Gaussian noise

$$\phi(x_{k-1}, v_{k-1}) = Dx_{k-1} + \Sigma_d$$

$$x_k \sim N(Dx_{k-1}, \Sigma_d)$$

State at time t comes from a transformation (D) of previous state ($k-1$) plus a noise term

- Observation model

- Measurement is the linearly transformed state plus Gaussian noise.

$$\psi(z_k, v_k) = Mx_k + \Sigma_m$$

$$z_k \sim N(Mx_k, \Sigma_m)$$

Measurement at time k comes from a transformation (M) of current state k plus a noise term

Linear Dynamics Models (LDM)



Notation

$$x_k \sim N(\mu, \Sigma)$$

- Random variable with Gaussian probability distribution that has mean vector μ and covariance matrix Σ
 - x and μ are d -dimensional, Σ is $d \times d$.
- For the unidimensional case, μ is a scalar and Σ is a 1×1 matrix \rightarrow the variance σ^2

$$x_t \sim N(\mu, \sigma^2)$$

$$(2\pi)^{-\frac{k}{2}} |\Sigma|^{-\frac{1}{2}} e^{-\frac{1}{2}(\mathbf{x}-\mu)' \Sigma^{-1}(\mathbf{x}-\mu)},$$



Linear Dynamics Models (LDM)

Linear dynamics models

- Example: linear velocity (1D points)
 - **State vector:** position p and velocity v

$$x_t = \begin{bmatrix} p_t \\ v_t \end{bmatrix} \quad \begin{aligned} p_t &= p_{t-1} + (\Delta t)v_{t-1} + \varepsilon \\ v_t &= v_{t-1} + \xi \end{aligned} \quad \begin{array}{l} \text{(greek letters} \\ \text{denote noise} \\ \text{terms)} \end{array}$$

$$x_t = D_t x_{t-1} + \text{noise} = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix} \begin{bmatrix} p_{t-1} \\ v_{t-1} \end{bmatrix} + \text{noise}$$

- **Measurement** is position only

$$z_t = M x_t + \text{noise} = \begin{bmatrix} 1 & 0 \end{bmatrix} \begin{bmatrix} p_t \\ v_t \end{bmatrix} + \text{noise}$$

$$\mathbf{x}_t \sim N(\mathbf{D}\mathbf{x}_{t-1}; \Sigma_d)$$

$$z_t \sim N(\mathbf{M}\mathbf{x}_t; \Sigma_m)$$

KALMAN FILTER

Tracking with LDM case: Kalman Filter

Kalman filter

- Optimal method for tracking **linear dynamical models** under the assumption of **Gaussian noise**
- The predicted/corrected state distributions are Gaussian
 - Only the mean and covariance should be estimated.
 - The calculations are easy (all the integrals can be done in closed form).

Tracking with LDM case: Kalman Filter


Kalman filter (1D)

Know corrected state from k-1,
and all measurements up to the
current one.

→ Predict distribution over next
state.

Know prediction of state, and
next measurement

→ Update distribution over
current state.



$$p(x_k | z_{1:k-1})$$

$$\mu_k^-, \sigma_k^-$$

Mean and std. dev. of predicted state:

$$p(x_k | z_{1:k})$$

$$\mu_k^+, \sigma_k^+$$

Mean and std. dev. of predicted state:

Tracking with LDM case: Kalman Filter

Kalman filter (1D): prediction

- Linear dynamic model defining predicted state evolution, with noise

$$x_k = N(d \cdot x_{k-1}, \sigma_d^2)$$

- Estimate the predicted distribution for next state

$$p(x_k | z_{0:k-1}) = N(\mu_k^-, (\sigma_k^-)^2)$$

- Update mean and variance

$$\mu_k^- = d \cdot \mu_{k-1}^+$$

$$(\sigma_k^-)^2 = \sigma_d^2 + (d \cdot \sigma_{k-1}^+)^2$$

Tracking with LDM case: Kalman Filter

Kalman filter (1D): correction

- Linear model defining the mapping of state to measurements:

$$z_k = N(m \cdot x_k, \sigma_m^2)$$

- Want to estimate corrected distribution given latest measurement:

$$P(x_k | z_{1:k}) = N(\mu_k^+, (\sigma_k^+)^2)$$

- Update mean & variance

$$\mu_k^+ = \frac{\mu_k^- \cdot \sigma_m^2 + m \cdot z_k \cdot (\sigma_k^-)^2}{\sigma_m^2 + m^2 \cdot (\sigma_k^-)^2}$$

$$(\sigma_k^+)^2 = \frac{\sigma_m^2 \cdot (\sigma_k^-)^2}{\sigma_m^2 + m^2 \cdot (\sigma_k^-)^2}$$

Tracking with LDM case: Kalman Filter

Kalman filter (1D): prediction vs. update

$$\mu_k^+ = \frac{\mu_k^- \cdot \sigma_m^2 + m \cdot z_k \cdot (\sigma_k^-)^2}{\sigma_m^2 + m^2 \cdot (\sigma_k^-)^2}$$

$$(\sigma_k^+)^2 = \frac{\sigma_m^2 \cdot (\sigma_k^-)^2}{\sigma_m^2 + m^2 \cdot (\sigma_k^-)^2}$$

- If there is no prediction uncertainty: $\sigma_k^- = 0$
 - **Measurement is ignored!**


$$\mu_k^+ = \mu_k^- \quad (\sigma_k^+)^2 = 0$$

- If there is no measurement uncertainty: $\sigma_m = 0$
 - **Prediction is ignored!**

$$\mu_k^+ = \frac{z_k}{m} \quad (\sigma_k^+)^2 = 0$$

Tracking with LDM case: Kalman Filter

Kalman filter (1D): example




State is 2D: position + velocity
Measurement is 1D: position

Slide credit: Kristen Grauman

Tracking with LDM case: Kalman Filter

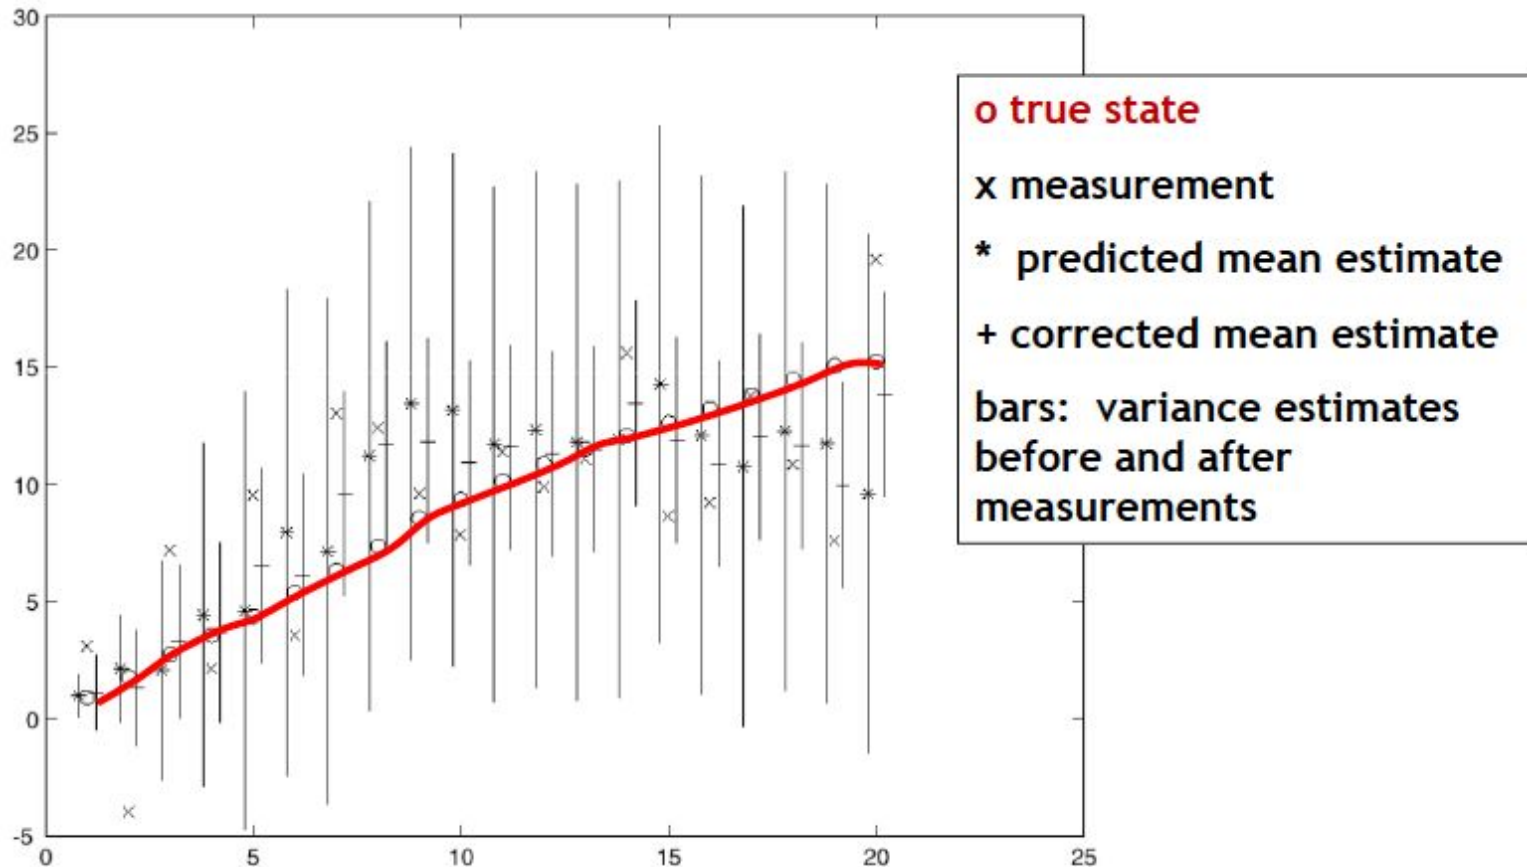
Kalman filter (1D): example



Slide credit: Kristen Grauman

Tracking with LDM case: Kalman Filter


Kalman filter (1D): example



Slide credit: Kristen Grauman

Tracking with LDM case: Kalman Filter


Kalman filter (1D): example



Slide credit: Kristen Grauman

Tracking with LDM case: Kalman Filter

Kalman filter (1D): example



Slide credit: Kristen Grauman

Tracking with LDM case: Kalman Filter

Kalman filter: generic case

- Vectors with more than one dimension:

Prediction

$$x_k^- = Dx_{k-1}^+$$

$$\Sigma_k^- = D\Sigma_k^+D^T + \Sigma_d$$

Update

$$x_k^+ = x_k^- + \boxed{K_k} \cdot \boxed{(z_k - M \cdot x_k^-)}$$

Kalman gain Residual

$$\Sigma_k^+ = (I - K_k M) \Sigma_k^-$$

$$\boxed{K_k = \Sigma_k^- M^T (M \Sigma_k^- M^T + \Sigma_m)^{-1}}$$

- More weight on residual when measurement error covariance approaches 0.
- Less weight on residual as a priori estimate error covariance approaches 0.

Slide credit: Kristen Grauman

Tracking with LDM case: Kalman Filter

Kalman filter


- Pros:
 - Gaussian densities everywhere
 - Simple updates, compact and efficient
 - Very established method, very well understood
 - Optimal for linear dynamic models under Gaussian noise
- Cons:
 - Unimodal distribution, only single hypothesis
 - Restricted class of motions defined by linear model

Modern visual trackers using Kalman Filter:

- Bewley, Alex et al. "Simple Online and Realtime Tracking." 2016 IEEE International Conference on Image Processing (ICIP)
- Nicolai Wojke et al. "Simple Online and Realtime Tracking with a Deep Association Metric", arXiv. 1703.07402, 2017

Tracking with LDM case: Kalman Filter

Kalman filter





Source: <https://github.com/Team-Neighborhood/Kalman-Filter-Image>

Tracking with LDM case: Kalman Filter

Kalman filter

- Linear model often does not describe accurately the dynamics of the problem
 - E.g. bouncing ball



Prediction is too far from true position to compensate



Slide credit: B. Leibe

Tracking with LDM case: Kalman Filter

Tracking issues

- Initialization
 - Often done manually
 - Background subtraction, detection can also be used
- Data association, multiple tracked objects
 - Occlusions, clutter
- Deformable and articulated objects
- Constructing accurate models of dynamics
 - E.g., Fitting parameters for a linear dynamics model
- Drift
 - Accumulation of errors over time

Tracking without LDM assumption

Other methods

- Extended Kalman filter
 - Removes linearity constraint on the state transition and observation models
- Mean-shift
 - Non-parametric technique
- Particle filter
 - Uses a sequential Montecarlo method
 - Multimodal distribution.
 - Removes linearity and gaussianity constraints